

```
In [1]: #Base Imports
import string
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
#Pre Processing Imports
import re
import nltk
from nltk.corpus import stopwords
nltk.download('stopwords')
nltk.download('wordnet')
from nltk.stem import WordNetLemmatizer
stop_words = set(stopwords.words("english"))
lemmatizer= WordNetLemmatizer()
from sklearn.model_selection import train_test_split
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
#Model Building Imports
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.models import Sequential
from tensorflow.keras.callbacks import EarlyStopping
from tensorflow.keras.layers import Dense, SimpleRNN, LSTM, Embedding, Dropout, GRU #Ac
#Model Evaluation imports
from sklearn import metrics
from sklearn.metrics import (classification_report, confusion_matrix ,
                             precision_recall_curve , precision_score, recall_score , accu
```

```
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data]   Package stopwords is already up-to-date!
[nltk_data] Downloading package wordnet to /root/nltk_data...
[nltk_data]   Package wordnet is already up-to-date!
```

```
In [2]: df = pd.read_csv('emotions_preprocessed.csv')
```

Data Preprocessing

```
In [3]: def lemmatization(text):
    lemmatizer= WordNetLemmatizer()

    text = text.split()

    text=[lemmatizer.lemmatize(y) for y in text]

    return " " .join(text)

def remove_stop_words(text):

    Text=[i for i in str(text).split() if i not in stop_words]
    return " " .join(Text)

def remove_numbers(text):
    text=''.join([i for i in text if not i.isdigit()])
    return text

def lower_case(text):
```

```

text = text.split()

text=[y.lower() for y in text]

return " " .join(text)

def remove_punctuations(text):
    ## Remove punctuations
    text = re.sub('[%s]' % re.escape("""!"#$%&'()*+,-./:;<=>?@[\\]^_`{|}~"""), ' ', text)
    text = text.replace(':', ',')

    ## remove extra whitespace
    text = re.sub('\s+', ' ', text)
    text = " ".join(text.split())
    return text.strip()

def remove_urls(text):
    url = re.compile(r'https?://\S+|www\.\S+')
    return url.sub('', text)

def clean(df):

    df.text=df.text.apply(lambda text : remove_stop_words(text))
    df.text=df.text.apply(lambda text : remove_numbers(text))
    df.text=df.text.apply(lambda text : remove_punctuations(text))
    df.text=df.text.apply(lambda text : remove_urls(text))
    df.text=df.text.apply(lambda text : lower_case(text))
    df.text=df.text.apply(lambda text : lemmatization(text))
    return df

```

In [4]: df = clean(df)

In [5]: X= df.text
y= df.labels
X_train, X_other, y_train, y_other = train_test_split(X, y, test_size=0.4, random_state=7)
X_val,X_test,y_val,y_test= train_test_split(X_other,y_other,test_size=0.5,random_state=7)
#60/20/20 split used

In [6]: X_train.head()

Out[6]: 31795 there's lot u lady here unfortunately 😞
6383 i surprised see name towing line show
20664 name worse he's thing plus took boi name
23929 probably would tell worry much wait bit try fe...
6505 you nailed it
Name: text, dtype: object

In [7]: y_train = le.fit_transform(y_train)
y_test = le.transform(y_test)
y_val = le.transform(y_val)

y_train = to_categorical(y_train)
y_test = to_categorical(y_test)
y_val = to_categorical(y_val)

In [8]: # Tokenize words
tokenizer = Tokenizer(oov_token='UNK')

```
tokenizer.fit_on_texts(X_train)
#tokenizer.fit_on_texts(pd.concat([X_train], axis=0))
```

In [9]: X_train

```
Out[9]: 31795          there's lot u lady here unfortunately 😞
6383          i surprised see name towing line show
20664          name worse he's thing plus took boi name
23929  probably would tell worry much wait bit try fe...
6505          you nailed it

...
21563  how supposed get link first place site trustwo...
25916          ooo i go far
44824          at job i can do get laid lot second one
21618          what word i m confused
23886          honestly that's wife material get that
Name: text, Length: 32396, dtype: object
```

```
In [10]: sequences_train = tokenizer.texts_to_sequences(X_train)
sequences_test = tokenizer.texts_to_sequences(X_test)
sequences_val = tokenizer.texts_to_sequences(X_val)
```

```
In [11]: max_len = max([len(t) for t in X_train])
max_len
```

Out[11]: 155

```
In [12]: X_train = pad_sequences(sequences_train, maxlen = max_len, truncating='pre')
X_test = pad_sequences(sequences_test, maxlen = max_len, truncating='pre')
X_val = pad_sequences(sequences_val, maxlen = max_len, truncating='pre')

vocabSize = len(tokenizer.index_word) + 1
print(f"Vocabulary size = {vocabSize}")

Vocabulary size = 20642
```

```
In [13]: np.unique(X_train)
```

Out[13]: array([0, 2, 3, ..., 20639, 20640, 20641], dtype=int32)

Text Representation Using Glove Embedding

```
In [14]: path_to_glove_file = 'glove.6B.300d.txt'
num_tokens = vocabSize
embedding_dim = 300
embeddings_index = {}
misses=0
hits=0
```

```
In [15]: with open(path_to_glove_file) as f:
    for line in f:
        word, coefs = line.split(maxsplit=1)
        coefs = np.fromstring(coefs, "f", sep=" ")
        embeddings_index[word] = coefs
    print("Found %s word vectors." % len(embeddings_index))

embedding_matrix = np.zeros((num_tokens, embedding_dim))
for word, i in tokenizer.word_index.items():
    embedding_vector = embeddings_index.get(word)
    if embedding_vector is not None:
        embedding_matrix[i] = embedding_vector
```

```

        hits += 1
    else:
        misses += 1
print("Converted %d words (%d misses)" % (hits, misses))

```

Found 8074 word vectors.
Converted 5106 words (15535 misses)

In [16]: `x_train.shape`

Out[16]: (32396, 155)

In [17]: `y_train.shape`

Out[17]: (32396, 14)

In [18]: *#use early stopping to control overfitting*
`callback = EarlyStopping(
 monitor="val_loss",
 patience=3,
 restore_best_weights=True,
)`

Simple RNN

In [19]: `rnn_model = Sequential()
rnn_model.add(Embedding(vocabSize, 300, input_length=X_train.shape[1], weights=[embeddin
rnn_model.add(SimpleRNN(units=64, return_sequences=True))
rnn_model.add(Dropout(0.5))
rnn_model.add(SimpleRNN(units=32))
rnn_model.add(Dropout(0.5))
rnn_model.add(Dense(14, activation='softmax'))
rnn_model.summary()`

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
embedding (Embedding)	(None, 155, 300)	6192600
simple_rnn (SimpleRNN)	(None, 155, 64)	23360
dropout (Dropout)	(None, 155, 64)	0
simple_rnn_1 (SimpleRNN)	(None, 32)	3104
dropout_1 (Dropout)	(None, 32)	0
dense (Dense)	(None, 14)	462
=====		
Total params: 6,219,526		
Trainable params: 26,926		
Non-trainable params: 6,192,600		

In [20]: `rnn_model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy`

In [21]: *# Fit model*
`history_rnn=rnn_model.fit(X_train,
 y_train,
 validation_data=(X_val, y_val),`

```

        verbose=1,
        batch_size=256,
        epochs=30,
        callbacks=[callback]
    )

```

```

Epoch 1/30
127/127 [=====] - 49s 306ms/step - loss: 2.6291 - accuracy: 0.2
060 - val_loss: 2.2951 - val_accuracy: 0.2997
Epoch 2/30
127/127 [=====] - 40s 315ms/step - loss: 2.3864 - accuracy: 0.2
828 - val_loss: 2.1779 - val_accuracy: 0.3385
Epoch 3/30
127/127 [=====] - 39s 304ms/step - loss: 2.2217 - accuracy: 0.3
418 - val_loss: 2.0948 - val_accuracy: 0.3589
Epoch 4/30
127/127 [=====] - 38s 299ms/step - loss: 2.1477 - accuracy: 0.3
633 - val_loss: 2.0946 - val_accuracy: 0.3596
Epoch 5/30
127/127 [=====] - 39s 309ms/step - loss: 2.1129 - accuracy: 0.3
725 - val_loss: 2.0740 - val_accuracy: 0.3717
Epoch 6/30
127/127 [=====] - 39s 305ms/step - loss: 2.0665 - accuracy: 0.3
860 - val_loss: 2.0155 - val_accuracy: 0.3946
Epoch 7/30
127/127 [=====] - 38s 299ms/step - loss: 2.0528 - accuracy: 0.3
939 - val_loss: 1.9979 - val_accuracy: 0.4032
Epoch 8/30
127/127 [=====] - 38s 301ms/step - loss: 2.0210 - accuracy: 0.4
030 - val_loss: 1.9862 - val_accuracy: 0.4084
Epoch 9/30
127/127 [=====] - 38s 299ms/step - loss: 1.9980 - accuracy: 0.4
148 - val_loss: 1.9640 - val_accuracy: 0.4124
Epoch 10/30
127/127 [=====] - 38s 302ms/step - loss: 1.9859 - accuracy: 0.4
146 - val_loss: 1.9599 - val_accuracy: 0.4156
Epoch 11/30
127/127 [=====] - 40s 312ms/step - loss: 1.9712 - accuracy: 0.4
201 - val_loss: 1.9569 - val_accuracy: 0.4191
Epoch 12/30
127/127 [=====] - 39s 309ms/step - loss: 1.9637 - accuracy: 0.4
239 - val_loss: 1.9496 - val_accuracy: 0.4160
Epoch 13/30
127/127 [=====] - 40s 319ms/step - loss: 1.9584 - accuracy: 0.4
250 - val_loss: 1.9402 - val_accuracy: 0.4194
Epoch 14/30
127/127 [=====] - 40s 317ms/step - loss: 2.1714 - accuracy: 0.3
585 - val_loss: 2.5909 - val_accuracy: 0.1667
Epoch 15/30
127/127 [=====] - 38s 301ms/step - loss: 2.1547 - accuracy: 0.3
527 - val_loss: 2.0482 - val_accuracy: 0.3835
Epoch 16/30
127/127 [=====] - 39s 310ms/step - loss: 2.0694 - accuracy: 0.3
870 - val_loss: 2.0175 - val_accuracy: 0.3926

```

```

In [22]: rnn_y_pred = rnn_model.predict(X_test)
rnn_y_pred_labels = np.argmax(rnn_y_pred, axis=1)
rnn_y_test_labels = np.argmax(y_test, axis=1)
print(classification_report(rnn_y_test_labels, rnn_y_pred_labels))

```

```

338/338 [=====] - 10s 28ms/step
              precision    recall  f1-score   support

     0           0.39         0.90         0.55         3147
     1           0.00         0.00         0.00          798
     2           0.00         0.00         0.00          367

```

3	0.33	0.05	0.08	243
4	0.00	0.00	0.00	688
5	0.47	0.56	0.51	1010
6	0.00	0.00	0.00	243
7	0.73	0.86	0.79	557
8	0.43	0.18	0.26	781
9	0.26	0.37	0.31	333
10	0.50	0.31	0.39	592
11	0.00	0.00	0.00	923
12	0.00	0.00	0.00	537
13	0.53	0.51	0.52	580
accuracy			0.43	10799
macro avg	0.26	0.27	0.24	10799
weighted avg	0.30	0.43	0.33	10799

```

/usr/local/lib/python3.9/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.9/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.9/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))

```

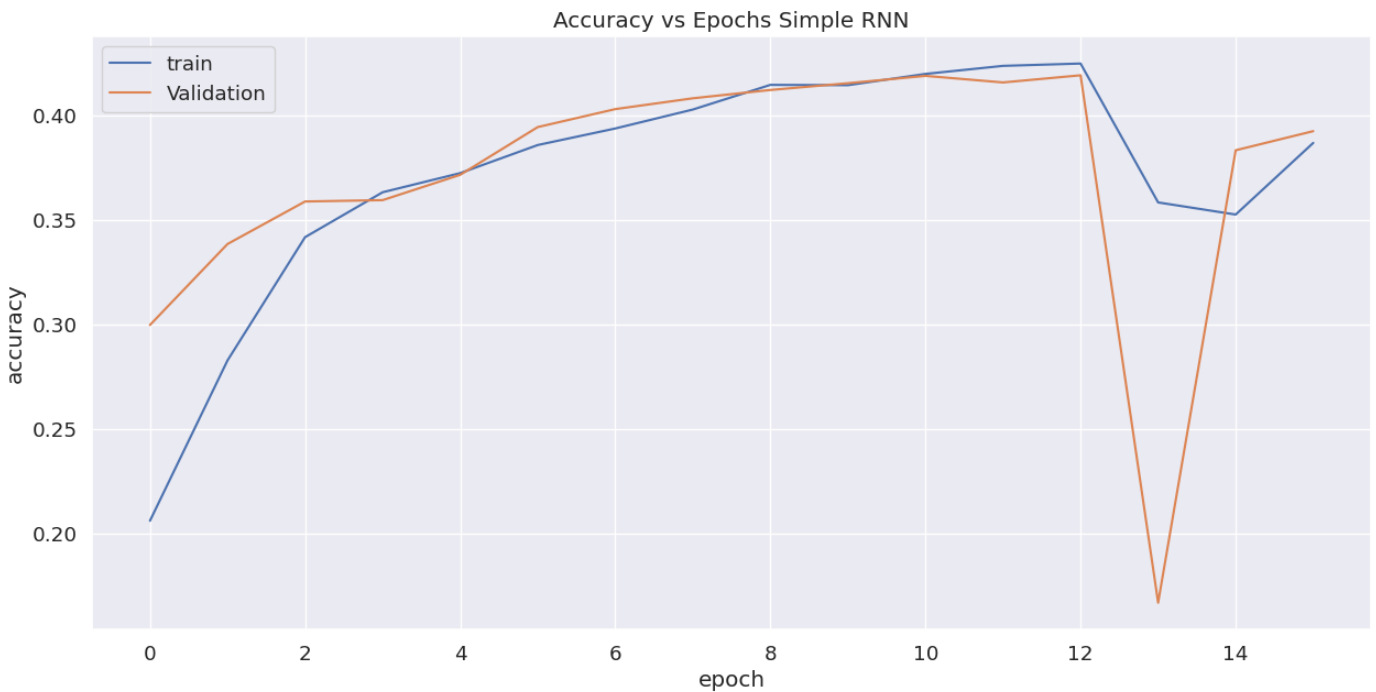
```

In [23]: conf_mat_rnn = confusion_matrix(rnn_y_test_labels, rnn_y_pred_labels)
plt.figure(figsize=(15,7))
# create a heatmap of the confusion matrix
sns.set(font_scale=1.2)
sns.heatmap(conf_mat_rnn, annot=True, fmt='d', cmap='Blues', cbar=False,
            xticklabels=['class 0', 'class 1', 'class 2', 'class 3', 'class 4', 'class 5',
                        'class 6', 'class 7', 'class 8', 'class 9', 'class 10',
                        'class 11', 'class 12', 'class 13'],
            yticklabels=['class 0', 'class 1', 'class 2', 'class 3', 'class 4', 'class 5',
                        'class 6', 'class 7', 'class 8', 'class 9', 'class 10',
                        'class 11', 'class 12', 'class 13'])
plt.xlabel('Predicted Labels', fontsize=14)
plt.ylabel('True Labels', fontsize=14)
plt.title('Confusion Matrix Simple RNN', fontsize=16)
plt.show()

```


True Labels \ Predicted Labels	class 0	class 1	class 2	class 3	class 4	class 5	class 6	class 7	class 8	class 9	class 10	class 11	class 12	class 13
class 0	2824	0	0	1	0	159	0	10	34	44	27	4	0	44
class 1	668	0	0	2	0	42	0	13	15	32	10	0	0	16
class 2	290	0	0	2	0	15	0	9	19	25	1	0	0	6
class 3	134	0	0	11	0	7	0	4	4	69	3	3	0	8
class 4	517	0	0	3	0	79	0	12	22	18	16	1	0	20
class 5	295	0	0	0	0	563	0	54	18	8	24	0	0	48
class 6	185	0	0	3	0	4	0	0	5	32	7	0	0	7
class 7	26	0	0	0	0	28	0	478	6	6	10	0	0	3
class 8	409	0	0	2	0	94	0	22	143	30	40	0	0	41
class 9	158	0	0	3	0	18	0	4	5	122	8	1	0	14
class 10	275	0	0	1	0	53	0	24	25	8	185	0	0	21
class 11	800	0	0	1	0	39	0	10	20	32	5	0	0	16
class 12	446	0	0	2	0	30	0	6	10	18	10	1	0	14
class 13	174	0	0	2	0	55	0	9	7	17	23	0	0	293

```
In [24]: plt.figure(figsize=(15,7))
plt.plot(history_rnn.history['accuracy'])
plt.plot(history_rnn.history['val_accuracy'])
plt.title('Accuracy vs Epochs Simple RNN')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'Validation'], loc='upper left')
plt.show()
```



```
In [25]: plt.figure(figsize=(15,7))
plt.plot(history_rnn.history['loss'])
plt.plot(history_rnn.history['val_loss'])
plt.title('Loss VS Epochs Simple RNN')
plt.ylabel('Loss')
plt.xlabel('epoch')
plt.legend(['train', 'Validation'], loc='upper left')
plt.show()
```

Loss VS Epochs Simple RNN



LSTM MODEL

```
In [26]: lstm_model = Sequential()
lstm_model.add(Embedding(vocabSize, 300, input_length=X_train.shape[1], weights=[embeddi
lstm_model.add(LSTM(units=64, return_sequences=True))
lstm_model.add(Dropout(0.5))
lstm_model.add(LSTM(units=32))
lstm_model.add(Dropout(0.5))
lstm_model.add(Dense(14, activation='softmax'))
lstm_model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	(None, 155, 300)	6192600
lstm (LSTM)	(None, 155, 64)	93440
dropout_2 (Dropout)	(None, 155, 64)	0
lstm_1 (LSTM)	(None, 32)	12416
dropout_3 (Dropout)	(None, 32)	0
dense_1 (Dense)	(None, 14)	462

```
=====  
Total params: 6,298,918  
Trainable params: 106,318  
Non-trainable params: 6,192,600  
=====
```

```
In [27]: lstm_model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accurac
```

```
In [28]: # Fit model
history_lstm=lstm_model.fit(X_train,
                             y_train,
                             validation_data=(X_val, y_val),
                             verbose=1,
```



```

        batch_size=256,
        epochs=30,
        callbacks=[callback]
    )

```

```

Epoch 1/30
127/127 [=====] - 9s 31ms/step - loss: 2.3150 - accuracy: 0.309
1 - val_loss: 2.1228 - val_accuracy: 0.3640
Epoch 2/30
127/127 [=====] - 3s 24ms/step - loss: 2.0699 - accuracy: 0.393
3 - val_loss: 1.9659 - val_accuracy: 0.4116
Epoch 3/30
127/127 [=====] - 3s 24ms/step - loss: 1.9727 - accuracy: 0.422
8 - val_loss: 1.9121 - val_accuracy: 0.4241
Epoch 4/30
127/127 [=====] - 3s 22ms/step - loss: 1.9249 - accuracy: 0.434
3 - val_loss: 1.8773 - val_accuracy: 0.4299
Epoch 5/30
127/127 [=====] - 3s 24ms/step - loss: 1.8906 - accuracy: 0.443
9 - val_loss: 1.8493 - val_accuracy: 0.4415
Epoch 6/30
127/127 [=====] - 3s 24ms/step - loss: 1.8615 - accuracy: 0.450
7 - val_loss: 1.8398 - val_accuracy: 0.4420
Epoch 7/30
127/127 [=====] - 3s 25ms/step - loss: 1.8404 - accuracy: 0.454
2 - val_loss: 1.8177 - val_accuracy: 0.4503
Epoch 8/30
127/127 [=====] - 3s 22ms/step - loss: 1.8192 - accuracy: 0.458
5 - val_loss: 1.8074 - val_accuracy: 0.4525
Epoch 9/30
127/127 [=====] - 3s 24ms/step - loss: 1.7998 - accuracy: 0.465
8 - val_loss: 1.7956 - val_accuracy: 0.4574
Epoch 10/30
127/127 [=====] - 3s 22ms/step - loss: 1.7881 - accuracy: 0.468
1 - val_loss: 1.8050 - val_accuracy: 0.4541
Epoch 11/30
127/127 [=====] - 3s 23ms/step - loss: 1.7740 - accuracy: 0.470
9 - val_loss: 1.7804 - val_accuracy: 0.4617
Epoch 12/30
127/127 [=====] - 3s 22ms/step - loss: 1.7537 - accuracy: 0.476
7 - val_loss: 1.7804 - val_accuracy: 0.4638
Epoch 13/30
127/127 [=====] - 3s 22ms/step - loss: 1.7424 - accuracy: 0.479
3 - val_loss: 1.7813 - val_accuracy: 0.4630
Epoch 14/30
127/127 [=====] - 3s 24ms/step - loss: 1.7239 - accuracy: 0.486
6 - val_loss: 1.7776 - val_accuracy: 0.4641
Epoch 15/30
127/127 [=====] - 3s 23ms/step - loss: 1.7140 - accuracy: 0.485
9 - val_loss: 1.7715 - val_accuracy: 0.4639
Epoch 16/30
127/127 [=====] - 3s 22ms/step - loss: 1.6969 - accuracy: 0.488
9 - val_loss: 1.7796 - val_accuracy: 0.4630
Epoch 17/30
127/127 [=====] - 3s 22ms/step - loss: 1.6865 - accuracy: 0.493
5 - val_loss: 1.7780 - val_accuracy: 0.4617
Epoch 18/30
127/127 [=====] - 3s 24ms/step - loss: 1.6736 - accuracy: 0.499
0 - val_loss: 1.7856 - val_accuracy: 0.4619

```

```

In [29]: lstm_y_pred = lstm_model.predict(X_test)
lstm_y_pred_labels = np.argmax(lstm_y_pred, axis=1)
lstm_y_test_labels = np.argmax(y_test, axis=1)
print(classification_report(lstm_y_test_labels, lstm_y_pred_labels))

```

```

338/338 [=====] - 3s 6ms/step

```

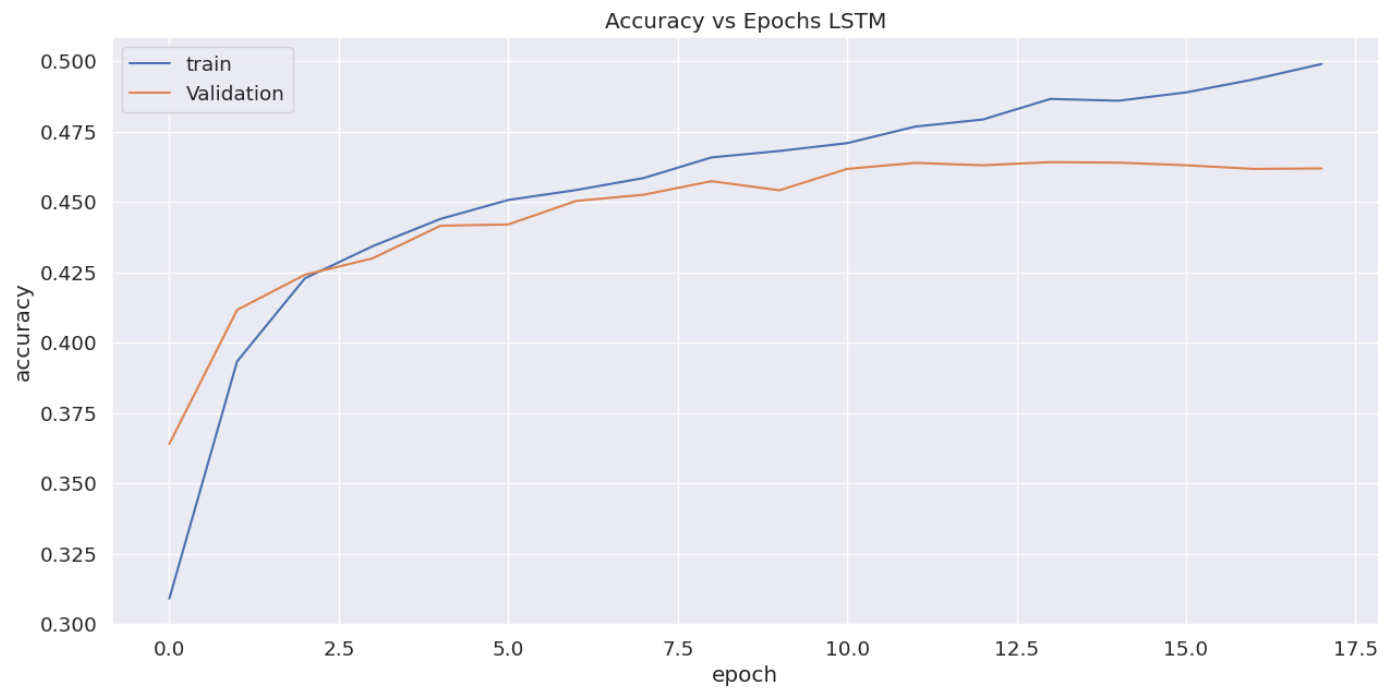
	precision	recall	f1-score	support
0	0.43	0.84	0.57	3147
1	0.44	0.21	0.29	798
2	0.48	0.22	0.30	367
3	0.49	0.25	0.33	243
4	0.53	0.09	0.15	688
5	0.59	0.55	0.57	1010
6	0.45	0.16	0.23	243
7	0.78	0.85	0.82	557
8	0.54	0.29	0.38	781
9	0.36	0.36	0.36	333
10	0.54	0.43	0.48	592
11	0.34	0.14	0.20	923
12	0.23	0.02	0.04	537
13	0.53	0.57	0.55	580
accuracy			0.48	10799
macro avg	0.48	0.35	0.38	10799
weighted avg	0.47	0.48	0.43	10799

```
In [30]: conf_mat_lstm = confusion_matrix(lstm_y_test_labels, lstm_y_pred_labels)
plt.figure(figsize=(15,7))
# create a heatmap of the confusion matrix
sns.set(font_scale=1.2)
sns.heatmap(conf_mat_lstm, annot=True, fmt='d', cmap='Blues', cbar=False,
            xticklabels=['class 0', 'class 1', 'class 2', 'class 3', 'class 4', 'class 5',
                        'class 6', 'class 7', 'class 8', 'class 9', 'class 10',
                        'class 11', 'class 12', 'class 13'],
            yticklabels=['class 0', 'class 1', 'class 2', 'class 3', 'class 4', 'class 5',
                        'class 6', 'class 7', 'class 8', 'class 9', 'class 10',
                        'class 11', 'class 12', 'class 13'])
plt.xlabel('Predicted Labels', fontsize=14)
plt.ylabel('True Labels', fontsize=14)
plt.title('Confusion Matrix LSTM', fontsize=16)
plt.show()
```

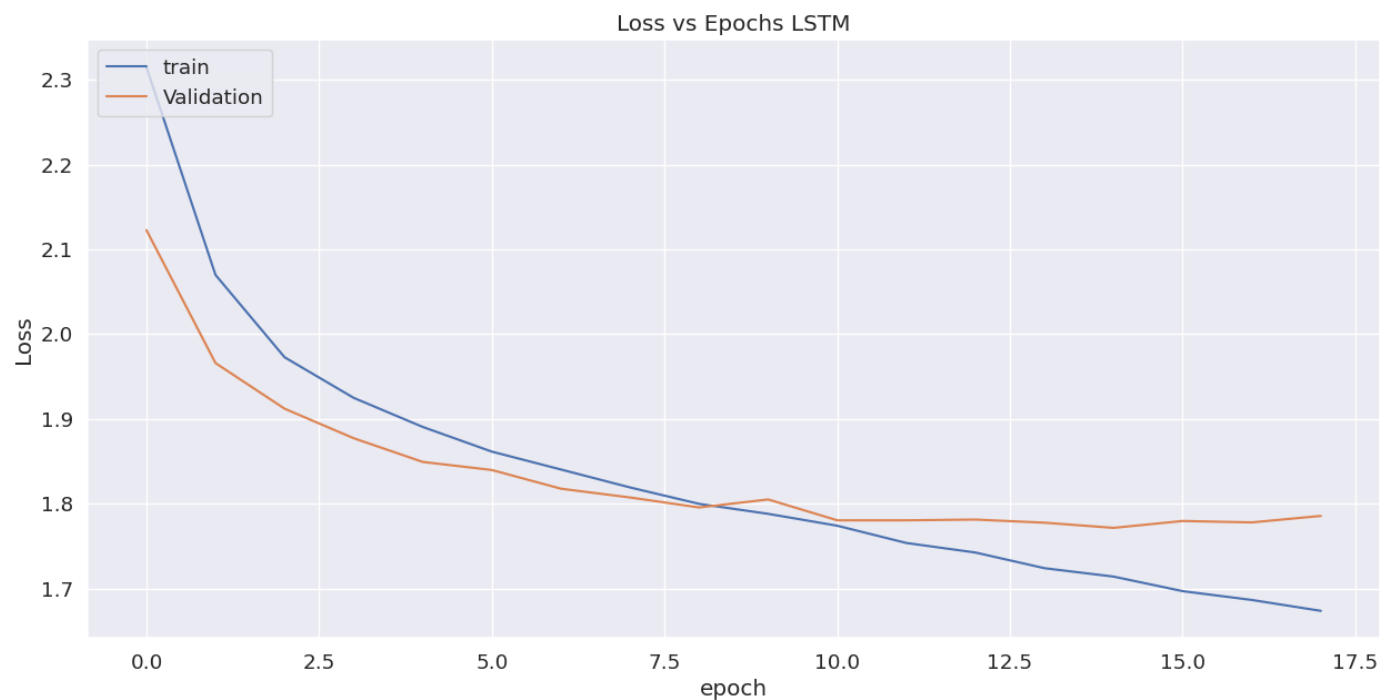
		Confusion Matrix LSTM													
True Labels	class 0	2638	90	18	12	13	93	7	5	33	36	44	84	11	63
	class 1	482	169	19	1	13	18	8	11	18	15	8	16	1	19
	class 2	219	8	79	0	3	10	3	0	10	14	7	10	1	3
	class 3	99	5	1	61	0	5	1	1	2	44	5	11	2	6
	class 4	427	12	13	8	62	63	2	15	18	12	15	9	6	26
	class 5	271	11	4	5	2	555	1	43	25	7	29	6	3	48
	class 6	131	10	4	6	0	6	38	0	3	13	6	23	1	2
	class 7	24	0	0	0	0	25	1	476	11	4	14	0	0	2
	class 8	340	20	9	10	4	45	5	20	225	12	30	18	2	41
	class 9	136	6	6	11	1	12	2	2	4	120	12	7	3	11
	class 10	211	11	5	0	6	29	1	15	22	5	252	6	2	27
	class 11	634	26	3	4	6	23	5	9	21	24	6	130	8	24
	class 12	371	7	2	3	7	18	4	6	10	21	11	49	12	16
	class 13	134	7	0	3	1	33	7	8	12	6	30	10	1	328
		class 0	class 1	class 2	class 3	class 4	class 5	class 6	class 7	class 8	class 9	class 10	class 11	class 12	class 13
		Predicted Labels													

```
In [31]: plt.figure(figsize=(15,7))
plt.plot(history_lstm.history['accuracy'])
plt.plot(history_lstm.history['val_accuracy'])
plt.title('Accuracy vs Epochs LSTM')
plt.ylabel('accuracy')
```

```
plt.xlabel('epoch')
plt.legend(['train', 'Validation'], loc='upper left')
plt.show()
```



```
In [32]: plt.figure(figsize=(15,7))
plt.plot(history_lstm.history['loss'])
plt.plot(history_lstm.history['val_loss'])
plt.title('Loss vs Epochs LSTM')
plt.ylabel('Loss')
plt.xlabel('epoch')
plt.legend(['train', 'Validation'], loc='upper left')
plt.show()
```



Gated Recurrent Unit (GRU)

```
In [33]: gru_model = Sequential()
gru_model.add(Embedding(vocabSize, 300, input_length=X_train.shape[1], weights=[embeddin
gru_model.add(GRU(units=64, return_sequences=True))
```

```

gru_model.add(Dropout(0.5))
gru_model.add(GRU(units=32))
gru_model.add(Dropout(0.5))
gru_model.add(Dense(14, activation='softmax'))
gru_model.summary()

```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
embedding_2 (Embedding)	(None, 155, 300)	6192600
gru (GRU)	(None, 155, 64)	70272
dropout_4 (Dropout)	(None, 155, 64)	0
gru_1 (GRU)	(None, 32)	9408
dropout_5 (Dropout)	(None, 32)	0
dense_2 (Dense)	(None, 14)	462
Total params: 6,272,742		
Trainable params: 80,142		
Non-trainable params: 6,192,600		

In [34]: `gru_model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy`

In [35]: `# Fit model`
`history_gru=gru_model.fit(X_train,`
`y_train,`
`validation_data=(X_val, y_val),`
`verbose=1,`
`batch_size=256,`
`epochs=30,`
`callbacks=[callback]`
`)`

```

Epoch 1/30
127/127 [=====] - 7s 27ms/step - loss: 2.2912 - accuracy: 0.323
8 - val_loss: 2.0435 - val_accuracy: 0.3900
Epoch 2/30
127/127 [=====] - 3s 24ms/step - loss: 2.0358 - accuracy: 0.403
0 - val_loss: 1.9425 - val_accuracy: 0.4211
Epoch 3/30
127/127 [=====] - 3s 23ms/step - loss: 1.9559 - accuracy: 0.426
4 - val_loss: 1.8913 - val_accuracy: 0.4349
Epoch 4/30
127/127 [=====] - 3s 21ms/step - loss: 1.9136 - accuracy: 0.441
4 - val_loss: 1.8558 - val_accuracy: 0.4448
Epoch 5/30
127/127 [=====] - 3s 21ms/step - loss: 1.8843 - accuracy: 0.448
1 - val_loss: 1.8324 - val_accuracy: 0.4526
Epoch 6/30
127/127 [=====] - 3s 23ms/step - loss: 1.8520 - accuracy: 0.452
3 - val_loss: 1.8050 - val_accuracy: 0.4554
Epoch 7/30
127/127 [=====] - 3s 21ms/step - loss: 1.8272 - accuracy: 0.459
1 - val_loss: 1.7996 - val_accuracy: 0.4567
Epoch 8/30
127/127 [=====] - 3s 23ms/step - loss: 1.8118 - accuracy: 0.463
1 - val_loss: 1.7818 - val_accuracy: 0.4618
Epoch 9/30
127/127 [=====] - 3s 21ms/step - loss: 1.7986 - accuracy: 0.467

```

```

3 - val_loss: 1.7753 - val_accuracy: 0.4624
Epoch 10/30
127/127 [=====] - 3s 23ms/step - loss: 1.7790 - accuracy: 0.472
1 - val_loss: 1.7660 - val_accuracy: 0.4687
Epoch 11/30
127/127 [=====] - 3s 22ms/step - loss: 1.7682 - accuracy: 0.475
1 - val_loss: 1.7601 - val_accuracy: 0.4698
Epoch 12/30
127/127 [=====] - 3s 23ms/step - loss: 1.7499 - accuracy: 0.479
9 - val_loss: 1.7649 - val_accuracy: 0.4664
Epoch 13/30
127/127 [=====] - 3s 23ms/step - loss: 1.7412 - accuracy: 0.481
6 - val_loss: 1.7620 - val_accuracy: 0.4669
Epoch 14/30
127/127 [=====] - 3s 23ms/step - loss: 1.7237 - accuracy: 0.485
2 - val_loss: 1.7526 - val_accuracy: 0.4696
Epoch 15/30
127/127 [=====] - 3s 24ms/step - loss: 1.7153 - accuracy: 0.486
9 - val_loss: 1.7490 - val_accuracy: 0.4729
Epoch 16/30
127/127 [=====] - 3s 21ms/step - loss: 1.7050 - accuracy: 0.488
0 - val_loss: 1.7485 - val_accuracy: 0.4746
Epoch 17/30
127/127 [=====] - 3s 21ms/step - loss: 1.6905 - accuracy: 0.490
5 - val_loss: 1.7522 - val_accuracy: 0.4708
Epoch 18/30
127/127 [=====] - 3s 23ms/step - loss: 1.6818 - accuracy: 0.493
8 - val_loss: 1.7602 - val_accuracy: 0.4702
Epoch 19/30
127/127 [=====] - 3s 24ms/step - loss: 1.6669 - accuracy: 0.497
9 - val_loss: 1.7660 - val_accuracy: 0.4696

```

```

In [36]: gru_y_pred = gru_model.predict(X_test)
gru_y_pred_labels = np.argmax(gru_y_pred, axis=1)
gru_y_test_labels = np.argmax(y_test, axis=1)
print(classification_report(gru_y_test_labels, gru_y_pred_labels))

```

```

338/338 [=====] - 2s 5ms/step

```

	precision	recall	f1-score	support
0	0.43	0.85	0.57	3147
1	0.45	0.23	0.31	798
2	0.49	0.19	0.27	367
3	0.53	0.28	0.36	243
4	0.55	0.12	0.19	688
5	0.58	0.57	0.58	1010
6	0.45	0.15	0.23	243
7	0.80	0.85	0.82	557
8	0.52	0.32	0.40	781
9	0.43	0.33	0.38	333
10	0.61	0.39	0.47	592
11	0.37	0.14	0.20	923
12	0.24	0.01	0.03	537
13	0.53	0.56	0.55	580
accuracy			0.48	10799
macro avg	0.50	0.36	0.38	10799
weighted avg	0.48	0.48	0.43	10799

```

In [37]: conf_mat_gru = confusion_matrix(gru_y_test_labels, gru_y_pred_labels)
plt.figure(figsize=(15,7))
# create a heatmap of the confusion matrix
sns.set(font_scale=1.2)
sns.heatmap(conf_mat_gru, annot=True, fmt='d', cmap='Blues', cbar=False,
            xticklabels=['class 0', 'class 1', 'class 2', 'class 3', 'class 4', 'class 5

```

```

        'class 6', 'class 7', 'class 8', 'class 9', 'class 10',
        'class 11', 'class 12', 'class 13'],
    yticklabels=['class 0', 'class 1', 'class 2', 'class 3', 'class 4', 'class 5',
        'class 6', 'class 7', 'class 8', 'class 9', 'class 10',
        'class 11', 'class 12', 'class 13'])
plt.xlabel('Predicted Labels', fontsize=14)
plt.ylabel('True Labels', fontsize=14)
plt.title('Confusion Matrix GRU', fontsize=16)
plt.show()

```


Confusion Matrix GRU

True Labels \ Predicted Labels	class 0	class 1	class 2	class 3	class 4	class 5	class 6	class 7	class 8	class 9	class 10	class 11	class 12	class 13
class 0	2665	97	14	8	14	109	9	4	43	18	29	80	5	52
class 1	465	185	16	6	13	23	10	11	23	12	3	12	2	17
class 2	231	12	68	3	4	9	1	2	12	11	2	7	0	5
class 3	117	5	1	67	0	6	1	1	2	31	3	3	1	5
class 4	438	11	10	5	80	53	1	14	22	10	10	8	2	24
class 5	261	17	3	2	8	575	1	36	25	4	20	6	3	49
class 6	139	8	4	5	1	5	37	0	6	9	7	18	2	2
class 7	25	0	0	1	1	29	0	472	13	3	9	0	0	4
class 8	338	15	8	8	2	51	6	18	252	9	14	18	0	42
class 9	138	4	6	10	2	16	2	1	9	111	11	10	2	11
class 10	233	6	3	0	8	24	2	12	34	6	228	4	1	31
class 11	632	37	4	3	2	27	4	9	21	16	5	127	7	29
class 12	388	11	3	4	5	23	4	5	8	12	6	44	8	16
class 13	148	2	0	4	5	33	5	7	14	7	25	3	1	326

```

In [38]: plt.figure(figsize=(15,7))
plt.plot(history_gru.history['accuracy'])
plt.plot(history_gru.history['val_accuracy'])
plt.title('Accuracy vs Epochs GRU')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'Validation'], loc='upper left')
plt.show()

```

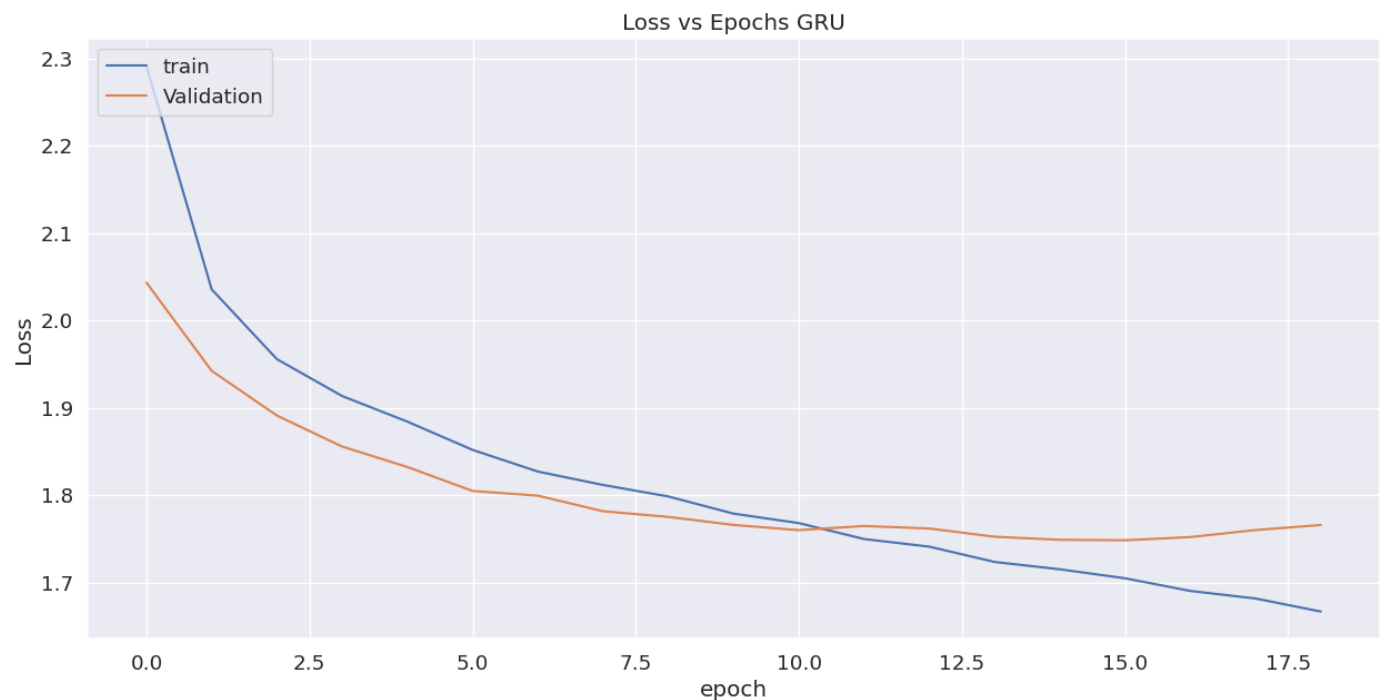


```

In [39]: plt.figure(figsize=(15,7))

```

```
plt.plot(history_gru.history['loss'])
plt.plot(history_gru.history['val_loss'])
plt.title('Loss vs Epochs GRU')
plt.ylabel('Loss')
plt.xlabel('epoch')
plt.legend(['train', 'Validation'], loc='upper left')
plt.show()
```



LSTM With GRU Layers

```
In [40]: lstmgru_model = Sequential()
lstmgru_model.add(Embedding(vocabSize, 300, input_length=X_train.shape[1], weights=[embe
lstmgru_model.add(LSTM(units=64, return_sequences=True))
lstmgru_model.add(Dropout(0.5))
lstmgru_model.add(GRU(units=32))
lstmgru_model.add(Dropout(0.5))
lstmgru_model.add(Dense(14, activation='softmax'))
lstmgru_model.summary()
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
embedding_3 (Embedding)	(None, 155, 300)	6192600
lstm_2 (LSTM)	(None, 155, 64)	93440
dropout_6 (Dropout)	(None, 155, 64)	0
gru_2 (GRU)	(None, 32)	9408
dropout_7 (Dropout)	(None, 32)	0
dense_3 (Dense)	(None, 14)	462
Total params: 6,295,910		
Trainable params: 103,310		
Non-trainable params: 6,192,600		

```
In [41]: lstmgru_model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accu
```

```
In [42]: # Fit model
history_lstmgru=lstmgru_model.fit(X_train,
                                   y_train,
                                   validation_data=(X_val, y_val),
                                   verbose=1,
                                   batch_size=256,
                                   epochs=30,
                                   callbacks=[callback]
                                   )
```

```
Epoch 1/30
127/127 [=====] - 8s 32ms/step - loss: 2.2739 - accuracy: 0.329
4 - val_loss: 2.0342 - val_accuracy: 0.3887
Epoch 2/30
127/127 [=====] - 3s 22ms/step - loss: 2.0292 - accuracy: 0.406
1 - val_loss: 1.9280 - val_accuracy: 0.4211
Epoch 3/30
127/127 [=====] - 3s 25ms/step - loss: 1.9461 - accuracy: 0.428
9 - val_loss: 1.8883 - val_accuracy: 0.4370
Epoch 4/30
127/127 [=====] - 3s 24ms/step - loss: 1.9063 - accuracy: 0.438
3 - val_loss: 1.8699 - val_accuracy: 0.4399
Epoch 5/30
127/127 [=====] - 3s 25ms/step - loss: 1.8798 - accuracy: 0.442
0 - val_loss: 1.8395 - val_accuracy: 0.4452
Epoch 6/30
127/127 [=====] - 3s 25ms/step - loss: 1.8534 - accuracy: 0.452
5 - val_loss: 1.8246 - val_accuracy: 0.4487
Epoch 7/30
127/127 [=====] - 3s 22ms/step - loss: 1.8298 - accuracy: 0.458
2 - val_loss: 1.8083 - val_accuracy: 0.4518
Epoch 8/30
127/127 [=====] - 3s 22ms/step - loss: 1.8126 - accuracy: 0.462
4 - val_loss: 1.7992 - val_accuracy: 0.4537
Epoch 9/30
127/127 [=====] - 3s 25ms/step - loss: 1.7887 - accuracy: 0.465
0 - val_loss: 1.8018 - val_accuracy: 0.4554
Epoch 10/30
127/127 [=====] - 3s 22ms/step - loss: 1.7739 - accuracy: 0.470
2 - val_loss: 1.7894 - val_accuracy: 0.4546
Epoch 11/30
127/127 [=====] - 3s 22ms/step - loss: 1.7559 - accuracy: 0.477
2 - val_loss: 1.7914 - val_accuracy: 0.4557
Epoch 12/30
127/127 [=====] - 3s 22ms/step - loss: 1.7430 - accuracy: 0.478
4 - val_loss: 1.7811 - val_accuracy: 0.4605
Epoch 13/30
127/127 [=====] - 3s 23ms/step - loss: 1.7289 - accuracy: 0.481
8 - val_loss: 1.7874 - val_accuracy: 0.4608
Epoch 14/30
127/127 [=====] - 3s 23ms/step - loss: 1.7181 - accuracy: 0.485
8 - val_loss: 1.7799 - val_accuracy: 0.4620
Epoch 15/30
127/127 [=====] - 3s 22ms/step - loss: 1.7021 - accuracy: 0.488
5 - val_loss: 1.7781 - val_accuracy: 0.4626
Epoch 16/30
127/127 [=====] - 3s 24ms/step - loss: 1.6842 - accuracy: 0.492
2 - val_loss: 1.7904 - val_accuracy: 0.4618
Epoch 17/30
127/127 [=====] - 3s 25ms/step - loss: 1.6740 - accuracy: 0.495
8 - val_loss: 1.7796 - val_accuracy: 0.4628
Epoch 18/30
```


127/127 [=====] - 3s 23ms/step - loss: 1.6545 - accuracy: 0.500
8 - val_loss: 1.7882 - val_accuracy: 0.4622

```
In [43]: lstmgru_y_pred = lstmgru_model.predict(X_test)
lstmgru_y_pred_labels = np.argmax(lstmgru_y_pred, axis=1)
lstmgru_y_test_labels = np.argmax(y_test, axis=1)
print(classification_report(lstmgru_y_test_labels, lstmgru_y_pred_labels))
```


```
338/338 [=====] - 2s 5ms/step
```

	precision	recall	f1-score	support
0	0.42	0.84	0.56	3147
1	0.42	0.25	0.31	798
2	0.50	0.19	0.27	367
3	0.52	0.24	0.33	243
4	0.51	0.11	0.18	688
5	0.61	0.55	0.58	1010
6	0.53	0.12	0.19	243
7	0.79	0.83	0.81	557
8	0.57	0.32	0.41	781
9	0.42	0.30	0.35	333
10	0.61	0.40	0.48	592
11	0.30	0.15	0.20	923
12	0.17	0.01	0.03	537
13	0.55	0.55	0.55	580
accuracy			0.48	10799
macro avg	0.50	0.35	0.38	10799
weighted avg	0.48	0.48	0.43	10799

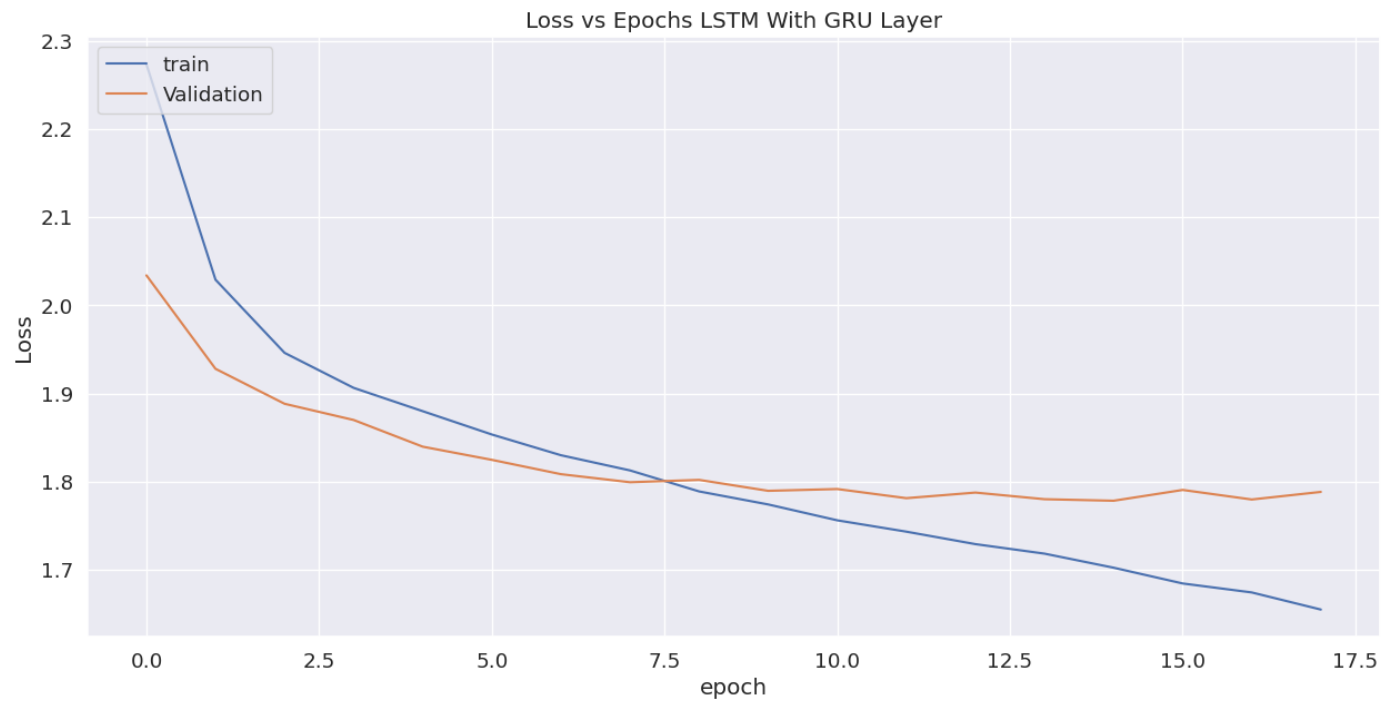
```
In [44]: conf_mat_lstmgru = confusion_matrix(lstmgru_y_test_labels, lstmgru_y_pred_labels)
plt.figure(figsize=(15,7))
# create a heatmap of the confusion matrix
sns.set(font_scale=1.2)
sns.heatmap(conf_mat_lstmgru, annot=True, fmt='d', cmap='Blues', cbar=False,
            xticklabels=['class 0', 'class 1', 'class 2', 'class 3', 'class 4', 'class 5',
                        'class 6', 'class 7', 'class 8', 'class 9', 'class 10',
                        'class 11', 'class 12', 'class 13'],
            yticklabels=['class 0', 'class 1', 'class 2', 'class 3', 'class 4', 'class 5',
                        'class 6', 'class 7', 'class 8', 'class 9', 'class 10',
                        'class 11', 'class 12', 'class 13'])
plt.xlabel('Predicted Labels', fontsize=14)
plt.ylabel('True Labels', fontsize=14)
plt.title('Confusion Matrix LSTM With GRU Layer', fontsize=16)
plt.show()
```

Confusion Matrix LSTM With GRU Layer														
True Labels	class 0	class 1	class 2	class 3	class 4	class 5	class 6	class 7	class 8	class 9	class 10	class 11	class 12	class 13
	2648	123	10	5	20	85	7	4	29	14	30	111	10	51
	462	199	15	3	11	19	7	12	17	10	4	19	4	16
	237	15	68	0	3	7	3	1	10	7	3	9	1	3
	118	5	1	59	0	6	1	2	1	31	4	9	0	6
	435	18	12	6	76	51	0	13	19	8	10	12	5	23
	289	16	4	3	7	551	0	39	19	8	18	9	1	46
	139	6	5	6	3	6	28	0	4	11	6	23	3	3
	28	4	0	0	1	27	0	465	15	3	11	0	0	3
	343	17	7	8	6	46	1	16	247	9	16	24	0	41
	148	5	7	12	1	10	2	2	4	99	11	18	3	11
	233	11	4	0	6	27	0	14	29	3	236	6	1	22
	627	45	0	7	5	20	1	9	21	17	4	136	8	23
	373	13	2	3	8	25	1	5	8	7	8	63	8	13
	161	2	0	2	1	25	2	8	12	6	25	13	2	321
Predicted Labels														

```
In [45]: plt.figure(figsize=(15,7))
plt.plot(history_lstmgru.history['accuracy'])
plt.plot(history_lstmgru.history['val_accuracy'])
plt.title('Accuracy vs Epochs LSTM With GRU Layer')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'Validation'], loc='upper left')
plt.show()
```



```
In [46]: plt.figure(figsize=(15,7))
plt.plot(history_lstmgru.history['loss'])
plt.plot(history_lstmgru.history['val_loss'])
plt.title('Loss vs Epochs LSTM With GRU Layer')
plt.ylabel('Loss')
plt.xlabel('epoch')
plt.legend(['train', 'Validation'], loc='upper left')
plt.show()
```



In [46]: