

```

In [1]: #Base Imports
import string
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
#Pre Processing Imports
import re
import nltk
from nltk.corpus import stopwords
nltk.download('stopwords')
nltk.download('wordnet')
from nltk.stem import WordNetLemmatizer
stop_words = set(stopwords.words("english"))
lemmatizer= WordNetLemmatizer()
from sklearn.model_selection import train_test_split
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
#Model Building Imports
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.models import Sequential
from tensorflow.keras.callbacks import EarlyStopping
from tensorflow.keras.layers import Dense, Embedding, Dropout, GRU #Activation, Flatten,
#Model Evaluation imports
from sklearn import metrics
from sklearn.metrics import (classification_report, confusion_matrix ,
                             precision_recall_curve ,precision_score,recall_score , accu

```

```

[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data]   Package stopwords is already up-to-date!
[nltk_data] Downloading package wordnet to /root/nltk_data...
[nltk_data]   Package wordnet is already up-to-date!

```

```

In [2]: df = pd.read_csv('emotions_preprocessed.csv')

```

```

In [3]: def lemmatization(text):
    lemmatizer= WordNetLemmatizer()

    text = text.split()

    text=[lemmatizer.lemmatize(y) for y in text]

    return " " .join(text)

def remove_stop_words(text):

    Text=[i for i in str(text).split() if i not in stop_words]
    return " " .join(Text)

def remove_numbers(text):
    text=''.join([i for i in text if not i.isdigit()])
    return text

def lower_case(text):

    text = text.split()

```

```

text=[y.lower() for y in text]

return " ".join(text)

def remove_punctuations(text):
    ## Remove punctuations
    text = re.sub('[%s]' % re.escape("""!"#$%&'()*+,-./:;<=>?@[\]^_`{|}~"""), ' ', text)
    text = text.replace(':', ',')

    ## remove extra whitespace
    text = re.sub('\s+', ' ', text)
    text = " ".join(text.split())
    return text.strip()

def remove_urls(text):
    url = re.compile(r'https?://\S+|www\.\S+')
    return url.sub('', text)

def clean(df):

    df.text=df.text.apply(lambda text : remove_stop_words(text))
    df.text=df.text.apply(lambda text : remove_numbers(text))
    df.text=df.text.apply(lambda text : remove_punctuations(text))
    df.text=df.text.apply(lambda text : remove_urls(text))
    df.text=df.text.apply(lambda text : lower_case(text))
    df.text=df.text.apply(lambda text : lemmatization(text))
    return df

```

```

In [4]: df = clean(df)
df

```

Out[4]:

	text	labels
0	my favourite food anything i cook myself	1
1	now himself everyone think he laugh screwing p...	1
2	why the fuck is bayless isoing	12
3	to make feel threatened	7
4	dirty southern wanker	12
...	...	...
53989	it s pretty dangerous state decides fictional ...	7
53990	i filed divorce morning hoping move next day so	11
53991	the last time happened i said no closed door	13
53992	i can't stand arrogant prick he's better thenf...	12
53993	but i like baby bang tiny voice	14

53994 rows x 2 columns

## 1ST SPLIT 90/5/5

```

In [5]: X= df.text

```

```

y= df.labels
X_train, X_other, y_train, y_other = train_test_split(X, y, test_size=0.1, random_state=
X_val,X_test,y_val,y_test= train_test_split(X_other,y_other,test_size=0.5,random_state=7

```

```

In [6]: le = LabelEncoder()
y_train = le.fit_transform(y_train)
y_test = le.transform(y_test)
y_val = le.transform(y_val)

y_train = to_categorical(y_train)
y_test = to_categorical(y_test)
y_val = to_categorical(y_val)

# Tokenize words
tokenizer = Tokenizer(oov_token='UNK')
tokenizer.fit_on_texts(pd.concat([X_train], axis=0))
sequences_train = tokenizer.texts_to_sequences(X_train)
sequences_test = tokenizer.texts_to_sequences(X_test)
sequences_val = tokenizer.texts_to_sequences(X_val)
max_len = max([len(t) for t in X_train])
X_train = pad_sequences(sequences_train, maxlen=max_len, truncating='pre')
X_test = pad_sequences(sequences_test, maxlen=max_len, truncating='pre')
X_val = pad_sequences(sequences_val, maxlen=max_len, truncating='pre')
vocabSize = len(tokenizer.index_word) + 1

#Text Representation Using Glove Embedding
path_to_glove_file = 'glove.6B.300d.txt'
num_tokens = vocabSize
embedding_dim = 300
embeddings_index = {}
misses=0
hits=0

with open(path_to_glove_file) as f:
    for line in f:
        word, coefs = line.split(maxsplit=1)
        coefs = np.fromstring(coefs, "f", sep=" ")
        embeddings_index[word] = coefs
print("Found %s word vectors." % len(embeddings_index))

embedding_matrix = np.zeros((num_tokens, embedding_dim))
for word, i in tokenizer.word_index.items():
    embedding_vector = embeddings_index.get(word)
    if embedding_vector is not None:
        embedding_matrix[i] = embedding_vector
        hits += 1
    else:
        misses += 1
print("Converted %d words (%d misses)" % (hits, misses))

#define early stopping to control overfitting
callback = EarlyStopping(
    monitor="val_loss",
    patience=3,
    restore_best_weights=True,
)

```

Found 37231 word vectors.  
Converted 13684 words (11651 misses)

```

In [7]: gru_model = Sequential()
gru_model.add(Embedding(vocabSize, 300, input_length=X_train.shape[1], weights=[embeddin
gru_model.add(GRU(units=64, return_sequences=True))

```

```

gru_model.add(Dropout(0.5))
gru_model.add(GRU(units=32))
gru_model.add(Dropout(0.5))
gru_model.add(Dense(14, activation='softmax'))
gru_model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 531, 300)	7600800
gru (GRU)	(None, 531, 64)	70272
dropout (Dropout)	(None, 531, 64)	0
gru_1 (GRU)	(None, 32)	9408
dropout_1 (Dropout)	(None, 32)	0
dense (Dense)	(None, 14)	462
Total params: 7,680,942		
Trainable params: 80,142		
Non-trainable params: 7,600,800		

In [8]: `gru_model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy`

In [9]: `# Fit model`  
`history_gru=gru_model.fit(X_train,`  
`y_train,`  
`validation_data=(X_val, y_val),`  
`verbose=1,`  
`batch_size=256,`  
`epochs=30,`  
`callbacks=[callback]`  
`)`

```

Epoch 1/30
190/190 [=====] - 21s 60ms/step - loss: 2.2238 - accuracy: 0.33
97 - val_loss: 1.8817 - val_accuracy: 0.4407
Epoch 2/30
190/190 [=====] - 11s 56ms/step - loss: 1.9140 - accuracy: 0.43
24 - val_loss: 1.7627 - val_accuracy: 0.4830
Epoch 3/30
190/190 [=====] - 11s 56ms/step - loss: 1.8275 - accuracy: 0.45
91 - val_loss: 1.6942 - val_accuracy: 0.4952
Epoch 4/30
190/190 [=====] - 11s 57ms/step - loss: 1.7689 - accuracy: 0.47
40 - val_loss: 1.6619 - val_accuracy: 0.5056
Epoch 5/30
190/190 [=====] - 11s 57ms/step - loss: 1.7325 - accuracy: 0.48
48 - val_loss: 1.6335 - val_accuracy: 0.5104
Epoch 6/30
190/190 [=====] - 11s 58ms/step - loss: 1.7006 - accuracy: 0.49
27 - val_loss: 1.6126 - val_accuracy: 0.5107
Epoch 7/30
190/190 [=====] - 11s 58ms/step - loss: 1.6774 - accuracy: 0.49
82 - val_loss: 1.6014 - val_accuracy: 0.5152
Epoch 8/30
190/190 [=====] - 11s 60ms/step - loss: 1.6532 - accuracy: 0.50
55 - val_loss: 1.5893 - val_accuracy: 0.5207
Epoch 9/30
190/190 [=====] - 12s 61ms/step - loss: 1.6384 - accuracy: 0.51

```



```

        'class 6', 'class 7', 'class 8', 'class 9', 'class 10',
        'class 11', 'class 12', 'class 13'])
plt.xlabel('Predicted Labels', fontsize=14)
plt.ylabel('True Labels', fontsize=14)
plt.title('Confusion Matrix 90/5/5 Split', fontsize=16)
plt.show()

```

Confusion Matrix 90/5/5 Split


class 0	606	16	8	3	10	29	5	3	13	8	7	35	7	13
class 1	105	39	9	1	2	2	1	3	5	5	1	11	6	5
class 2	37	4	32	0	1	5	1	0	3	2	3	1	1	0
class 3	18	1	0	20	0	1	0	0	0	8	1	1	0	3
class 4	100	1	5	1	18	17	0	3	6	0	2	5	2	8
class 5	49	2	3	0	3	176	2	8	10	2	6	2	0	15
class 6	10	1	2	0	0	1	28	0	3	1	0	3	3	0
class 7	5	0	0	1	0	14	0	109	3	0	2	0	0	0
class 8	59	3	5	0	0	14	3	8	78	2	9	8	0	17
class 9	36	1	1	3	2	2	2	0	0	23	6	11	2	2
class 10	46	0	6	0	0	9	0	1	12	0	63	3	0	4
class 11	106	3	1	2	2	9	7	3	4	4	4	97	5	5
class 12	71	2	1	1	1	2	3	3	4	6	1	24	13	4
class 13	28	1	1	1	1	8	1	1	3	4	7	4	1	77

Predicted Labels

```

In [12]: plt.figure(figsize=(15,7))
plt.plot(history_gru.history['accuracy'])
plt.plot(history_gru.history['val_accuracy'])
plt.title('Accuracy vs Epochs 90/5/5 Split')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'Validation'], loc='upper left')
plt.show()

```

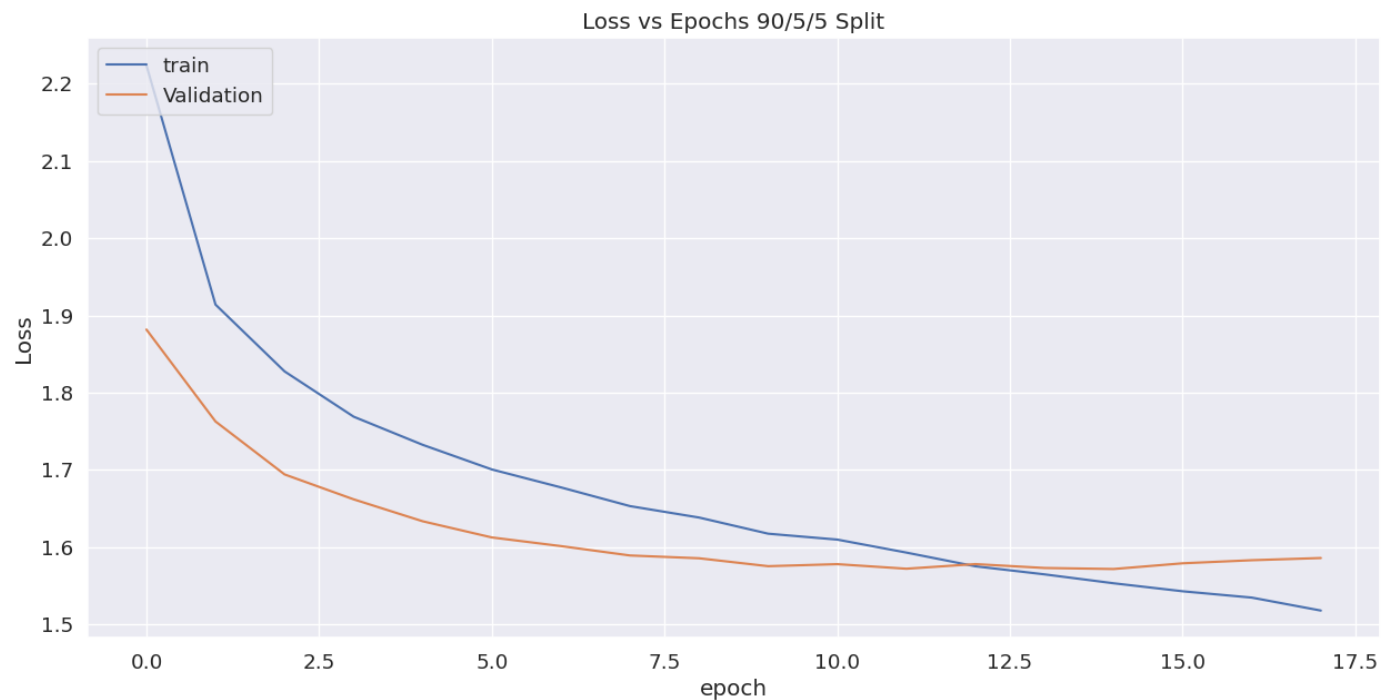


```

In [13]: plt.figure(figsize=(15,7))
plt.plot(history_gru.history['loss'])
plt.plot(history_gru.history['val_loss'])
plt.title('Loss vs Epochs 90/5/5 Split')

```

```
plt.ylabel('Loss')
plt.xlabel('epoch')
plt.legend(['train', 'Validation'], loc='upper left')
plt.show()
```



## 2ND SPLIT 10/45/45

```
In [14]: X= df.text
y= df.labels
X_train, X_other, y_train, y_other = train_test_split(X, y, test_size=0.9, random_state=
X_val,X_test,y_val,y_test= train_test_split(X_other,y_other,test_size=0.5,random_state=7
```

```
In [15]: le = LabelEncoder()
y_train = le.fit_transform(y_train)
y_test = le.transform(y_test)
y_val = le.transform(y_val)

y_train = to_categorical(y_train)
y_test = to_categorical(y_test)
y_val = to_categorical(y_val)

# Tokenize words
tokenizer = Tokenizer(oov_token='UNK')
tokenizer.fit_on_texts(pd.concat([X_train], axis=0))
sequences_train = tokenizer.texts_to_sequences(X_train)
sequences_test = tokenizer.texts_to_sequences(X_test)
sequences_val = tokenizer.texts_to_sequences(X_val)
max_len = max([len(t) for t in X_train])
X_train = pad_sequences(sequences_train, maxlen=max_len, truncating='pre')
X_test = pad_sequences(sequences_test, maxlen=max_len, truncating='pre')
X_val = pad_sequences(sequences_val, maxlen=max_len, truncating='pre')
vocabSize = len(tokenizer.index_word) + 1

#Text Representation Using Glove Embedding
path_to_glove_file = 'glove.6B.300d.txt'
num_tokens = vocabSize
embedding_dim = 300
embeddings_index = {}
misses=0
hits=0
```





```
        verbose=1,  
        batch_size=256,  
        epochs=30,  
        callbacks=[callback]  
    )
```

```
Epoch 1/30  
22/22 [=====] - 6s 87ms/step - loss: 2.4794 - accuracy: 0.2493  
- val_loss: 2.3368 - val_accuracy: 0.3081  
Epoch 2/30  
22/22 [=====] - 2s 78ms/step - loss: 2.3570 - accuracy: 0.2949  
- val_loss: 2.2740 - val_accuracy: 0.3264  
Epoch 3/30  
22/22 [=====] - 1s 53ms/step - loss: 2.2868 - accuracy: 0.3239  
- val_loss: 2.2061 - val_accuracy: 0.3490  
Epoch 4/30  
22/22 [=====] - 2s 80ms/step - loss: 2.2142 - accuracy: 0.3449  
- val_loss: 2.1364 - val_accuracy: 0.3639  
Epoch 5/30  
22/22 [=====] - 1s 50ms/step - loss: 2.1558 - accuracy: 0.3571  
- val_loss: 2.0823 - val_accuracy: 0.3780  
Epoch 6/30  
22/22 [=====] - 1s 50ms/step - loss: 2.0834 - accuracy: 0.3780  
- val_loss: 2.0186 - val_accuracy: 0.4003  
Epoch 7/30  
22/22 [=====] - 1s 50ms/step - loss: 2.0293 - accuracy: 0.3991  
- val_loss: 1.9718 - val_accuracy: 0.4116  
Epoch 8/30  
22/22 [=====] - 2s 77ms/step - loss: 1.9875 - accuracy: 0.4067  
- val_loss: 1.9521 - val_accuracy: 0.4166  
Epoch 9/30  
22/22 [=====] - 1s 51ms/step - loss: 1.9456 - accuracy: 0.4225  
- val_loss: 1.9108 - val_accuracy: 0.4338  
Epoch 10/30  
22/22 [=====] - 2s 78ms/step - loss: 1.9091 - accuracy: 0.4382  
- val_loss: 1.8975 - val_accuracy: 0.4378  
Epoch 11/30  
22/22 [=====] - 2s 78ms/step - loss: 1.8758 - accuracy: 0.4395  
- val_loss: 1.8878 - val_accuracy: 0.4405  
Epoch 12/30  
22/22 [=====] - 2s 89ms/step - loss: 1.8507 - accuracy: 0.4519  
- val_loss: 1.8760 - val_accuracy: 0.4453  
Epoch 13/30  
22/22 [=====] - 1s 57ms/step - loss: 1.8215 - accuracy: 0.4580  
- val_loss: 1.8672 - val_accuracy: 0.4487  
Epoch 14/30  
22/22 [=====] - 1s 51ms/step - loss: 1.7969 - accuracy: 0.4619  
- val_loss: 1.8670 - val_accuracy: 0.4497  
Epoch 15/30  
22/22 [=====] - 1s 50ms/step - loss: 1.7848 - accuracy: 0.4745  
- val_loss: 1.8521 - val_accuracy: 0.4525  
Epoch 16/30  
22/22 [=====] - 2s 78ms/step - loss: 1.7413 - accuracy: 0.4764  
- val_loss: 1.8530 - val_accuracy: 0.4532  
Epoch 17/30  
22/22 [=====] - 1s 52ms/step - loss: 1.7263 - accuracy: 0.4862  
- val_loss: 1.8679 - val_accuracy: 0.4546  
Epoch 18/30  
22/22 [=====] - 2s 78ms/step - loss: 1.7154 - accuracy: 0.4866  
- val_loss: 1.8508 - val_accuracy: 0.4555  
Epoch 19/30  
22/22 [=====] - 1s 56ms/step - loss: 1.6827 - accuracy: 0.4879  
- val_loss: 1.8429 - val_accuracy: 0.4593  
Epoch 20/30  
22/22 [=====] - 1s 55ms/step - loss: 1.6665 - accuracy: 0.4981  
- val_loss: 1.8631 - val_accuracy: 0.4578
```

```
Epoch 21/30
22/22 [=====] - 1s 51ms/step - loss: 1.6361 - accuracy: 0.5044
- val_loss: 1.8546 - val_accuracy: 0.4524
Epoch 22/30
22/22 [=====] - 2s 80ms/step - loss: 1.6085 - accuracy: 0.5116
- val_loss: 1.8519 - val_accuracy: 0.4539
```

```
In [18]: gru_y_pred = gru_model.predict(X_test)
gru_y_pred_labels = np.argmax(gru_y_pred, axis=1)
gru_y_test_labels = np.argmax(y_test, axis=1)
print(classification_report(gru_y_test_labels, gru_y_pred_labels))
```


```
760/760 [=====] - 4s 5ms/step
```

	precision	recall	f1-score	support
0	0.43	0.83	0.57	7134
1	0.34	0.01	0.02	1683
2	0.25	0.00	0.00	813
3	0.42	0.03	0.05	512
4	0.13	0.00	0.01	1516
5	0.52	0.63	0.57	2338
6	0.30	0.10	0.15	476
7	0.75	0.80	0.77	1206
8	0.44	0.29	0.35	1799
9	0.36	0.27	0.31	834
10	0.51	0.31	0.38	1264
11	0.36	0.39	0.38	2035
12	0.16	0.02	0.03	1330
13	0.54	0.45	0.49	1358
accuracy			0.46	24298
macro avg	0.39	0.30	0.29	24298
weighted avg	0.41	0.46	0.38	24298

```
In [19]: conf_mat_gru = confusion_matrix(gru_y_test_labels, gru_y_pred_labels)
plt.figure(figsize=(15,7))
# create a heatmap of the confusion matrix
sns.set(font_scale=1.2)
sns.heatmap(conf_mat_gru, annot=True, fmt='d', cmap='Blues', cbar=False,
            xticklabels=['class 0', 'class 1', 'class 2', 'class 3', 'class 4', 'class 5',
                        'class 6', 'class 7', 'class 8', 'class 9', 'class 10',
                        'class 11', 'class 12', 'class 13'],
            yticklabels=['class 0', 'class 1', 'class 2', 'class 3', 'class 4', 'class 5',
                        'class 6', 'class 7', 'class 8', 'class 9', 'class 10',
                        'class 11', 'class 12', 'class 13'])
plt.xlabel('Predicted Labels', fontsize=14)
plt.ylabel('True Labels', fontsize=14)
plt.title('Confusion Matrix 10/40/40 Split', fontsize=16)
plt.show()
```


Confusion Matrix 10/40/40 Split														
True Labels	class 0	class 1	class 2	class 3	class 4	class 5	class 6	class 7	class 8	class 9	class 10	class 11	class 12	class 13
	5955	6	1	2	4	305	12	23	144	46	56	445	24	111
	1311	19	1	1	4	70	12	27	44	30	23	101	10	30
	591	5	2	1	8	52	11	2	29	16	29	46	10	11
	242	1	0	13	1	23	23	7	17	92	9	67	4	13
	1048	3	0	0	4	143	5	33	69	30	39	74	10	58
	501	3	0	1	3	1483	1	73	79	1	47	46	11	89
	210	4	1	1	0	9	49	3	18	51	5	100	11	14
	76	0	0	0	1	84	1	962	31	22	11	4	0	14
	770	3	1	0	0	181	3	59	529	32	52	98	3	68
	369	2	0	6	1	50	10	5	28	225	29	86	14	9
	523	0	0	2	1	147	2	31	79	9	389	30	0	51
	983	3	0	0	3	74	14	19	54	29	15	795	13	33
	830	7	2	0	0	76	16	12	51	23	17	251	21	24
	452	0	0	4	1	134	6	21	28	18	40	39	3	612
Predicted Labels														

```
In [20]: plt.figure(figsize=(15,7))
plt.plot(history_gru.history['accuracy'])
plt.plot(history_gru.history['val_accuracy'])
plt.title('Accuracy vs Epochs 10/40/40 Split')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'Validation'], loc='upper left')
plt.show()
```



```
In [21]: plt.figure(figsize=(15,7))
plt.plot(history_gru.history['loss'])
plt.plot(history_gru.history['val_loss'])
plt.title('Loss vs Epochs 10/40/40 Split')
plt.ylabel('Loss')
plt.xlabel('epoch')
plt.legend(['train', 'Validation'], loc='upper left')
plt.show()
```

Loss vs Epochs 10/40/40 Split



## 3RD SPLIT 25/25/50

```
In [22]: X= df.text
y= df.labels
X_train, X_other, y_train, y_other = train_test_split(X, y, test_size=0.75, random_state=
X_val,X_test,y_val,y_test= train_test_split(X_other,y_other,test_size=0.33,random_state=
```

```
In [23]: le = LabelEncoder()
y_train = le.fit_transform(y_train)
y_test = le.transform(y_test)
y_val = le.transform(y_val)

y_train = to_categorical(y_train)
y_test = to_categorical(y_test)
y_val = to_categorical(y_val)

# Tokenize words
tokenizer = Tokenizer(oov_token='UNK')
tokenizer.fit_on_texts(pd.concat([X_train], axis=0))
sequences_train = tokenizer.texts_to_sequences(X_train)
sequences_test = tokenizer.texts_to_sequences(X_test)
sequences_val = tokenizer.texts_to_sequences(X_val)
max_len = max([len(t) for t in X_train])
X_train = pad_sequences(sequences_train, maxlen=max_len, truncating='pre')
X_test = pad_sequences(sequences_test, maxlen=max_len, truncating='pre')
X_val = pad_sequences(sequences_val, maxlen=max_len, truncating='pre')
vocabSize = len(tokenizer.index_word) + 1

#Text Representation Using Glove Embedding
path_to_glove_file = 'glove.6B.300d.txt'
num_tokens = vocabSize
embedding_dim = 300
embeddings_index = {}
misses=0
hits=0

with open(path_to_glove_file) as f:
    for line in f:
        word, coefs = line.split(maxsplit=1)
```

```

        coefs = np.fromstring(coefs, "f", sep=" ")
        embeddings_index[word] = coefs
print("Found %s word vectors." % len(embeddings_index))

embedding_matrix = np.zeros((num_tokens, embedding_dim))
for word, i in tokenizer.word_index.items():
    embedding_vector = embeddings_index.get(word)
    if embedding_vector is not None:
        embedding_matrix[i] = embedding_vector
        hits += 1
    else:
        misses += 1
print("Converted %d words (%d misses)" % (hits, misses))

#define early stopping to control overfitting
callback = EarlyStopping(
    monitor="val_loss",
    patience=3,
    restore_best_weights=True,
)

```

Found 65215 word vectors.  
Converted 9880 words (3161 misses)

```

In [24]: gru_model = Sequential()
gru_model.add(Embedding(vocabSize, 300, input_length=X_train.shape[1], weights=[embeddin
gru_model.add(GRU(units=64, return_sequences=True))
gru_model.add(Dropout(0.5))
gru_model.add(GRU(units=32))
gru_model.add(Dropout(0.5))
gru_model.add(Dense(14, activation='softmax'))
gru_model.summary()
gru_model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy

```

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
embedding_2 (Embedding)	(None, 145, 300)	3912600
gru_4 (GRU)	(None, 145, 64)	70272
dropout_4 (Dropout)	(None, 145, 64)	0
gru_5 (GRU)	(None, 32)	9408
dropout_5 (Dropout)	(None, 32)	0
dense_2 (Dense)	(None, 14)	462

=====  
 Total params: 3,992,742  
 Trainable params: 80,142  
 Non-trainable params: 3,912,600  
 =====

```

In [25]: # Fit model
history_gru=gru_model.fit(X_train,
                           y_train,
                           validation_data=(X_val, y_val),
                           verbose=1,
                           batch_size=256,
                           epochs=30,
                           callbacks=[callback]
)

```

```

Epoch 1/30
53/53 [=====] - 6s 48ms/step - loss: 2.4337 - accuracy: 0.2614
- val_loss: 2.2699 - val_accuracy: 0.3338
Epoch 2/30
53/53 [=====] - 2s 32ms/step - loss: 2.2260 - accuracy: 0.3412
- val_loss: 2.0427 - val_accuracy: 0.3881
Epoch 3/30
53/53 [=====] - 2s 42ms/step - loss: 2.0578 - accuracy: 0.3868
- val_loss: 1.9234 - val_accuracy: 0.4231
Epoch 4/30
53/53 [=====] - 2s 43ms/step - loss: 1.9803 - accuracy: 0.4136
- val_loss: 1.8669 - val_accuracy: 0.4399
Epoch 5/30
53/53 [=====] - 2s 42ms/step - loss: 1.9113 - accuracy: 0.4327
- val_loss: 1.8265 - val_accuracy: 0.4543
Epoch 6/30
53/53 [=====] - 2s 32ms/step - loss: 1.8709 - accuracy: 0.4500
- val_loss: 1.8102 - val_accuracy: 0.4587
Epoch 7/30
53/53 [=====] - 2s 33ms/step - loss: 1.8305 - accuracy: 0.4580
- val_loss: 1.7662 - val_accuracy: 0.4705
Epoch 8/30
53/53 [=====] - 2s 42ms/step - loss: 1.8009 - accuracy: 0.4695
- val_loss: 1.7462 - val_accuracy: 0.4759
Epoch 9/30
53/53 [=====] - 2s 34ms/step - loss: 1.7579 - accuracy: 0.4804
- val_loss: 1.7385 - val_accuracy: 0.4800
Epoch 10/30
53/53 [=====] - 2s 44ms/step - loss: 1.7389 - accuracy: 0.4801
- val_loss: 1.7224 - val_accuracy: 0.4815
Epoch 11/30
53/53 [=====] - 2s 33ms/step - loss: 1.7089 - accuracy: 0.4901
- val_loss: 1.7201 - val_accuracy: 0.4846
Epoch 12/30
53/53 [=====] - 2s 33ms/step - loss: 1.6862 - accuracy: 0.4973
- val_loss: 1.7129 - val_accuracy: 0.4864
Epoch 13/30
53/53 [=====] - 2s 33ms/step - loss: 1.6593 - accuracy: 0.5000
- val_loss: 1.7059 - val_accuracy: 0.4864
Epoch 14/30
53/53 [=====] - 2s 42ms/step - loss: 1.6410 - accuracy: 0.5106
- val_loss: 1.7275 - val_accuracy: 0.4858
Epoch 15/30
53/53 [=====] - 2s 33ms/step - loss: 1.6233 - accuracy: 0.5128
- val_loss: 1.7008 - val_accuracy: 0.4895
Epoch 16/30
53/53 [=====] - 2s 34ms/step - loss: 1.5957 - accuracy: 0.5192
- val_loss: 1.7043 - val_accuracy: 0.4898
Epoch 17/30
53/53 [=====] - 2s 34ms/step - loss: 1.5802 - accuracy: 0.5244
- val_loss: 1.7120 - val_accuracy: 0.4898
Epoch 18/30
53/53 [=====] - 2s 42ms/step - loss: 1.5532 - accuracy: 0.5301
- val_loss: 1.7133 - val_accuracy: 0.4850

```

```

In [26]: gru_y_pred = gru_model.predict(X_test)
gru_y_pred_labels = np.argmax(gru_y_pred, axis=1)
gru_y_test_labels = np.argmax(y_test, axis=1)
print(classification_report(gru_y_test_labels, gru_y_pred_labels))

```

```

418/418 [=====] - 3s 5ms/step
              precision    recall  f1-score   support

     0       0.45         0.86         0.59         3953
     1       0.41         0.04         0.08          940
     2       0.55         0.14         0.22          461

```


3	0.61	0.18	0.28	302
4	0.28	0.06	0.10	794
5	0.59	0.65	0.61	1218
6	0.44	0.23	0.30	268
7	0.78	0.80	0.79	691
8	0.54	0.34	0.42	1010
9	0.42	0.32	0.37	459
10	0.50	0.37	0.42	696
11	0.45	0.35	0.40	1125
12	0.22	0.01	0.02	701
13	0.57	0.50	0.53	746

accuracy			0.49	13364
macro avg	0.49	0.35	0.37	13364
weighted avg	0.48	0.49	0.43	13364


```
In [27]: conf_mat_gru = confusion_matrix(gru_y_test_labels, gru_y_pred_labels)
plt.figure(figsize=(15,7))
# create a heatmap of the confusion matrix
sns.set(font_scale=1.2)
sns.heatmap(conf_mat_gru, annot=True, fmt='d', cmap='Blues', cbar=False,
            xticklabels=['class 0', 'class 1', 'class 2', 'class 3', 'class 4', 'class 5',
                        'class 6', 'class 7', 'class 8', 'class 9', 'class 10',
                        'class 11', 'class 12', 'class 13'],
            yticklabels=['class 0', 'class 1', 'class 2', 'class 3', 'class 4', 'class 5',
                        'class 6', 'class 7', 'class 8', 'class 9', 'class 10',
                        'class 11', 'class 12', 'class 13'])
plt.xlabel('Predicted Labels', fontsize=14)
plt.ylabel('True Labels', fontsize=14)
plt.title('Confusion Matrix 25/25/50 Split', fontsize=16)
plt.show()
```

Confusion Matrix 25/25/50 Split														
True Labels	class 0	class 1	class 2	class 3	class 4	class 5	class 6	class 7	class 8	class 9	class 10	class 11	class 12	class 13
	3395	13	14	6	35	126	14	9	58	31	40	147	5	60
	711	41	9	3	13	24	11	20	25	16	14	29	5	19
	277	14	65	2	5	34	8	3	15	7	12	13	3	3
	143	1	0	54	1	6	3	3	7	45	7	22	0	10
	533	5	3	3	49	67	2	16	29	4	26	24	4	29
	244	1	11	0	22	787	1	38	23	5	22	14	1	49
	118	5	2	6	3	4	62	0	7	17	2	34	4	4
	33	1	0	0	0	40	1	551	27	5	21	1	0	11
	416	3	3	2	15	66	7	31	347	9	32	37	1	41
	212	1	3	7	3	22	2	1	9	149	13	31	1	5
	269	1	6	1	7	58	1	10	37	7	256	13	0	30
	577	7	0	2	9	24	16	13	20	26	14	399	3	15
	483	5	2	2	7	27	8	7	16	19	12	96	8	9
	206	3	1	1	4	60	5	5	19	11	38	17	2	374
Predicted Labels														

```
In [28]: plt.figure(figsize=(15,7))
plt.plot(history_gru.history['accuracy'])
plt.plot(history_gru.history['val_accuracy'])
plt.title('Accuracy vs Epochs 25/25/50 Split')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'Validation'], loc='upper left')
plt.show()
```



```
In [29]: plt.figure(figsize=(15,7))
plt.plot(history_gru.history['loss'])
plt.plot(history_gru.history['val_loss'])
plt.title('Loss vs Epochs 25/25/50 Split')
plt.ylabel('Loss')
plt.xlabel('epoch')
plt.legend(['train', 'Validation'], loc='upper left')
plt.show()
```



## 4TH SPLIT 25/50/25

```
In [30]: X= df.text
y= df.labels
X_train, X_other, y_train, y_other = train_test_split(X, y, test_size=0.75, random_state=
X_val,X_test,y_val,y_test= train_test_split(X_other,y_other,test_size=0.67,random_state=
```

```
In [31]: le =LabelEncoder()
```



```

y_train = le.fit_transform(y_train)
y_test = le.transform(y_test)
y_val = le.transform(y_val)

y_train = to_categorical(y_train)
y_test = to_categorical(y_test)
y_val = to_categorical(y_val)

# Tokenize words
tokenizer = Tokenizer(oov_token='UNK')
tokenizer.fit_on_texts(pd.concat([X_train], axis=0))
sequences_train = tokenizer.texts_to_sequences(X_train)
sequences_test = tokenizer.texts_to_sequences(X_test)
sequences_val = tokenizer.texts_to_sequences(X_val)
max_len = max([len(t) for t in X_train])
X_train = pad_sequences(sequences_train, maxlen=max_len, truncating='pre')
X_test = pad_sequences(sequences_test, maxlen=max_len, truncating='pre')
X_val = pad_sequences(sequences_val, maxlen=max_len, truncating='pre')
vocabSize = len(tokenizer.index_word) + 1

#Text Representation Using Glove Embedding
path_to_glove_file = 'glove.6B.300d.txt'
num_tokens = vocabSize
embedding_dim = 300
embeddings_index = {}
misses=0
hits=0

with open(path_to_glove_file) as f:
    for line in f:
        word, coefs = line.split(maxsplit=1)
        coefs = np.fromstring(coefs, "f", sep=" ")
        embeddings_index[word] = coefs
print("Found %s word vectors." % len(embeddings_index))

embedding_matrix = np.zeros((num_tokens, embedding_dim))
for word, i in tokenizer.word_index.items():
    embedding_vector = embeddings_index.get(word)
    if embedding_vector is not None:
        embedding_matrix[i] = embedding_vector
        hits += 1
    else:
        misses += 1
print("Converted %d words (%d misses)" % (hits, misses))

#define early stopping to control overfitting
callback = EarlyStopping(
    monitor="val_loss",
    patience=3,
    restore_best_weights=True,
)

```

Found 69272 word vectors.  
 Converted 9970 words (3071 misses)

```

In [32]: gru_model = Sequential()
gru_model.add(Embedding(vocabSize, 300, input_length=X_train.shape[1], weights=[embeddin
gru_model.add(GRU(units=64, return_sequences=True))
gru_model.add(Dropout(0.5))
gru_model.add(GRU(units=32))
gru_model.add(Dropout(0.5))
gru_model.add(Dense(14, activation='softmax'))
gru_model.summary()
gru_model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy

```

Model: "sequential\_3"

Layer (type)	Output Shape	Param #
embedding_3 (Embedding)	(None, 145, 300)	3912600
gru_6 (GRU)	(None, 145, 64)	70272
dropout_6 (Dropout)	(None, 145, 64)	0
gru_7 (GRU)	(None, 32)	9408
dropout_7 (Dropout)	(None, 32)	0
dense_3 (Dense)	(None, 14)	462

=====  
Total params: 3,992,742  
Trainable params: 80,142  
Non-trainable params: 3,912,600  
=====

```
In [33]: # Fit model
history_gru=gru_model.fit(X_train,
                           y_train,
                           validation_data=(X_val, y_val),
                           verbose=1,
                           batch_size=256,
                           epochs=30,
                           callbacks=[callback]
                           )
```

```
Epoch 1/30
53/53 [=====] - 6s 45ms/step - loss: 2.4354 - accuracy: 0.2721
- val_loss: 2.2576 - val_accuracy: 0.3222
Epoch 2/30
53/53 [=====] - 1s 25ms/step - loss: 2.2312 - accuracy: 0.3353
- val_loss: 2.0779 - val_accuracy: 0.3623
Epoch 3/30
53/53 [=====] - 2s 30ms/step - loss: 2.0886 - accuracy: 0.3777
- val_loss: 1.9588 - val_accuracy: 0.4100
Epoch 4/30
53/53 [=====] - 1s 27ms/step - loss: 1.9850 - accuracy: 0.4132
- val_loss: 1.8693 - val_accuracy: 0.4432
Epoch 5/30
53/53 [=====] - 1s 26ms/step - loss: 1.9139 - accuracy: 0.4347
- val_loss: 1.8175 - val_accuracy: 0.4607
Epoch 6/30
53/53 [=====] - 1s 25ms/step - loss: 1.8628 - accuracy: 0.4495
- val_loss: 1.7857 - val_accuracy: 0.4682
Epoch 7/30
53/53 [=====] - 2s 30ms/step - loss: 1.8301 - accuracy: 0.4612
- val_loss: 1.7681 - val_accuracy: 0.4726
Epoch 8/30
53/53 [=====] - 1s 25ms/step - loss: 1.7834 - accuracy: 0.4696
- val_loss: 1.7476 - val_accuracy: 0.4739
Epoch 9/30
53/53 [=====] - 2s 30ms/step - loss: 1.7671 - accuracy: 0.4747
- val_loss: 1.7354 - val_accuracy: 0.4771
Epoch 10/30
53/53 [=====] - 2s 30ms/step - loss: 1.7407 - accuracy: 0.4790
- val_loss: 1.7237 - val_accuracy: 0.4804
Epoch 11/30
53/53 [=====] - 2s 30ms/step - loss: 1.7130 - accuracy: 0.4911
- val_loss: 1.7150 - val_accuracy: 0.4846
```

```

Epoch 12/30
53/53 [=====] - 2s 30ms/step - loss: 1.6855 - accuracy: 0.4951
- val_loss: 1.7113 - val_accuracy: 0.4835
Epoch 13/30
53/53 [=====] - 2s 31ms/step - loss: 1.6634 - accuracy: 0.4980
- val_loss: 1.7086 - val_accuracy: 0.4863
Epoch 14/30
53/53 [=====] - 1s 26ms/step - loss: 1.6429 - accuracy: 0.5070
- val_loss: 1.7076 - val_accuracy: 0.4860
Epoch 15/30
53/53 [=====] - 1s 25ms/step - loss: 1.6293 - accuracy: 0.5087
- val_loss: 1.7023 - val_accuracy: 0.4883
Epoch 16/30
53/53 [=====] - 1s 25ms/step - loss: 1.6111 - accuracy: 0.5121
- val_loss: 1.7070 - val_accuracy: 0.4873
Epoch 17/30
53/53 [=====] - 1s 25ms/step - loss: 1.5717 - accuracy: 0.5251
- val_loss: 1.7156 - val_accuracy: 0.4884
Epoch 18/30
53/53 [=====] - 1s 26ms/step - loss: 1.5629 - accuracy: 0.5281
- val_loss: 1.7172 - val_accuracy: 0.4890

```

```

In [34]: gru_y_pred = gru_model.predict(X_test)
gru_y_pred_labels = np.argmax(gru_y_pred, axis=1)
gru_y_test_labels = np.argmax(y_test, axis=1)
print(classification_report(gru_y_test_labels, gru_y_pred_labels))

```

```

848/848 [=====] - 5s 6ms/step
              precision    recall  f1-score   support

     0         0.46         0.81         0.59         8047
     1         0.44         0.07         0.12         1879
     2         0.47         0.15         0.23          908
     3         0.46         0.20         0.28          581
     4         0.32         0.01         0.02         1635
     5         0.61         0.62         0.62         2597
     6         0.47         0.30         0.37          562
     7         0.77         0.83         0.80         1382
     8         0.51         0.38         0.43         2015
     9         0.34         0.35         0.34          871
    10         0.51         0.39         0.44         1418
    11         0.40         0.44         0.42         2286
    12         0.36         0.03         0.05         1455
    13         0.51         0.56         0.53         1497

 accuracy         0.49         27133
 macro avg        0.47         0.37         0.37         27133
 weighted avg     0.48         0.49         0.44         27133

```

```


In [35]: conf_mat_gru = confusion_matrix(gru_y_test_labels, gru_y_pred_labels)
plt.figure(figsize=(15,7))
# create a heatmap of the confusion matrix
sns.set(font_scale=1.2)
sns.heatmap(conf_mat_gru, annot=True, fmt='d', cmap='Blues', cbar=False,
            xticklabels=['class 0', 'class 1', 'class 2', 'class 3', 'class 4', 'class 5',
                        'class 6', 'class 7', 'class 8', 'class 9', 'class 10',
                        'class 11', 'class 12', 'class 13'],
            yticklabels=['class 0', 'class 1', 'class 2', 'class 3', 'class 4', 'class 5',
                        'class 6', 'class 7', 'class 8', 'class 9', 'class 10',
                        'class 11', 'class 12', 'class 13'])
plt.xlabel('Predicted Labels', fontsize=14)
plt.ylabel('True Labels', fontsize=14)
plt.title('Confusion Matrix 25/50/25 Split', fontsize=16)
plt.show()

```

Confusion Matrix 25/50/25 Split

True Labels	class 0	class 1	class 2	class 3	class 4	class 5	class 6	class 7	class 8	class 9	class 10	class 11	class 12	class 13
class 0	6489	50	40	24	13	266	27	20	171	87	103	552	14	191
class 1	1247	133	28	9	4	59	29	31	47	48	38	140	5	61
class 2	476	28	137	6	3	61	8	1	32	40	41	66	0	9
class 3	217	5	0	119	2	9	15	1	15	118	13	39	0	28
class 4	1086	14	13	8	21	131	6	36	65	34	54	77	12	78
class 5	495	2	22	5	11	1618	3	104	97	9	52	53	5	121
class 6	190	8	5	10	0	7	171	1	6	52	12	75	5	20
class 7	55	1	1	1	0	64	2	1142	55	13	24	5	0	19
class 8	759	10	13	11	0	118	10	51	758	29	59	85	3	109
class 9	335	7	3	31	1	25	15	4	20	303	26	76	7	18
class 10	472	9	18	7	3	101	4	30	96	15	547	39	2	75
class 11	950	17	1	9	3	48	32	25	45	59	19	1016	12	50
class 12	869	11	6	11	5	67	23	13	38	63	23	250	37	39
class 13	342	9	2	8	0	86	21	18	37	25	66	45	0	838
	class 0	class 1	class 2	class 3	class 4	class 5	class 6	class 7	class 8	class 9	class 10	class 11	class 12	class 13
	Predicted Labels													

```
In [36]: plt.figure(figsize=(15,7))
plt.plot(history_gru.history['accuracy'])
plt.plot(history_gru.history['val_accuracy'])
plt.title('Accuracy vs Epochs 25/25/50 Split')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'Validation'], loc='upper left')
plt.show()
```



```
In [37]: plt.figure(figsize=(15,7))
plt.plot(history_gru.history['loss'])
plt.plot(history_gru.history['val_loss'])
plt.title('Loss vs Epochs 25/50/25 Split')
plt.ylabel('Loss')
plt.xlabel('epoch')
plt.legend(['train', 'Validation'], loc='upper left')
plt.show()
```

