

```
In [1]: #Base Imports
import string
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
#Pre Processing Imports
import re
import nltk
from nltk.corpus import stopwords
nltk.download('stopwords')
nltk.download('wordnet')
from nltk.stem import WordNetLemmatizer , SnowballStemmer
stop_words = set(stopwords.words("english"))
lemmatizer= WordNetLemmatizer()
stemmer = SnowballStemmer("english")
from sklearn.model_selection import train_test_split
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
#Model Building Imports
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.models import Sequential
from tensorflow.keras.callbacks import EarlyStopping
from tensorflow.keras.layers import Dense, SimpleRNN , LSTM, Embedding, Dropout, GRU , Bi
#Model Evaluation imports
from sklearn import metrics
from sklearn.metrics import (classification_report, confusion_matrix ,
                             precision_recall_curve , precision_score, recall_score , accu
```

```
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data]   Package stopwords is already up-to-date!
[nltk_data] Downloading package wordnet to /root/nltk_data...
[nltk_data]   Package wordnet is already up-to-date!
```

```
In [2]: df = pd.read_csv('emotions_preprocessed.csv')
```

```
In [3]: def stemming(text):
    text = text.split()
    text = [stemmer.stem(y) for y in text]

    return " ".join(text)

def lemmatization(text):
    text = text.split()
    text=[lemmatizer.lemmatize(y) for y in text]

    return " " .join(text)

def remove_stop_words(text):

    Text=[i for i in str(text).split() if i not in stop_words]
    return " ".join(Text)

def remove_numbers(text):
    text=''.join([i for i in text if not i.isdigit()])
    return text
```

```

def lower_case(text):

    text = text.split()

    text=[y.lower() for y in text]

    return " ".join(text)

def remove_punctuations(text):
    ## Remove punctuations
    text = re.sub('[%s]' % re.escape("'\"!\"#$%&'()*+,-./:;<=>?@[\\]^_`{|}~\""), ' ', text)
    text = text.replace(':', ',')

    ## remove extra whitespace
    text = re.sub('\s+', ' ', text)
    text = " ".join(text.split())
    return text.strip()

def remove_urls(text):
    url = re.compile(r'https?://\S+|www\.\S+')
    return url.sub('', text)

def clean_stopwords_removed_lemma(df):

    df.text=df.text.apply(lambda text : remove_stop_words(text))
    df.text=df.text.apply(lambda text : remove_numbers(text))
    df.text=df.text.apply(lambda text : remove_punctuations(text))
    df.text=df.text.apply(lambda text : remove_urls(text))
    df.text=df.text.apply(lambda text : lower_case(text))
    df.text=df.text.apply(lambda text : lemmatization(text))
    return df

def clean_stopwords_removed_stemming(df):

    df.text=df.text.apply(lambda text : remove_stop_words(text))
    df.text=df.text.apply(lambda text : remove_numbers(text))
    df.text=df.text.apply(lambda text : remove_punctuations(text))
    df.text=df.text.apply(lambda text : remove_urls(text))
    df.text=df.text.apply(lambda text : lower_case(text))
    df.text=df.text.apply(lambda text : stemming(text))
    return df

def clean_stopwords_present_stemming(df):

    df.text=df.text.apply(lambda text : remove_numbers(text))
    df.text=df.text.apply(lambda text : remove_punctuations(text))
    df.text=df.text.apply(lambda text : remove_urls(text))
    df.text=df.text.apply(lambda text : lower_case(text))
    df.text=df.text.apply(lambda text : stemming(text))
    return df

def clean_stopwords_present_lemma(df):
    df.text=df.text.apply(lambda text : remove_numbers(text))
    df.text=df.text.apply(lambda text : remove_punctuations(text))
    df.text=df.text.apply(lambda text : remove_urls(text))
    df.text=df.text.apply(lambda text : lower_case(text))
    df.text=df.text.apply(lambda text : lemmatization(text))
    return df

```

In [4]: df_clean_stopwords_removed_stemming = clean_stopwords_removed_stemming(df)

```
df_clean_stopwords_present_stemming = clean_stopwords_present_stemming(df)
df_clean_stopwords_present_lemma = clean_stopwords_present_lemma(df)
```

Stemming With Stop Words Removed

```
In [5]: X= df_clean_stopwords_removed_stemming.text
y= df_clean_stopwords_removed_stemming.labels
X_train, X_other, y_train, y_other = train_test_split(X, y, test_size=0.1, random_state=
X_val,X_test,y_val,y_test= train_test_split(X_other,y_other,test_size=0.5,random_state=7
```

```
In [6]: le = LabelEncoder()
y_train = le.fit_transform(y_train)
y_test = le.transform(y_test)
y_val = le.transform(y_val)

y_train = to_categorical(y_train)
y_test = to_categorical(y_test)
y_val = to_categorical(y_val)
# Tokenize words
tokenizer = Tokenizer(oov_token='UNK')
tokenizer.fit_on_texts(pd.concat([X_train], axis=0))
sequences_train = tokenizer.texts_to_sequences(X_train)
sequences_test = tokenizer.texts_to_sequences(X_test)
sequences_val = tokenizer.texts_to_sequences(X_val)
max_len = max([len(t) for t in X_train])
X_train = pad_sequences(sequences_train, maxlen=max_len, truncating='pre')
X_test = pad_sequences(sequences_test, maxlen=max_len, truncating='pre')
X_val = pad_sequences(sequences_val, maxlen=max_len, truncating='pre')
vocabSize = len(tokenizer.index_word) + 1

#Text Representation Using Glove Embedding
path_to_glove_file = 'glove.6B.300d.txt'
num_tokens = vocabSize
embedding_dim = 300
embeddings_index = {}
misses=0
hits=0
with open(path_to_glove_file) as f:
    for line in f:
        word, coefs = line.split(maxsplit=1)
        coefs = np.fromstring(coefs, "f", sep=" ")
        embeddings_index[word] = coefs
print("Found %s word vectors." % len(embeddings_index))

embedding_matrix = np.zeros((num_tokens, embedding_dim))
for word, i in tokenizer.word_index.items():
    embedding_vector = embeddings_index.get(word)
    if embedding_vector is not None:
        embedding_matrix[i] = embedding_vector
        hits += 1
    else:
        misses += 1
print("Converted %d words (%d misses)" % (hits, misses))

#define early stopping to control overfitting
callback = EarlyStopping(
    monitor="val_loss",
    patience=3,
    restore_best_weights=True,
)
```

Found 23858 word vectors.

```
In [7]: gru_model = Sequential()
gru_model.add(Embedding(vocabSize, 300, input_length=X_train.shape[1], weights=[embeddin
gru_model.add(GRU(units=64, return_sequences=True))
gru_model.add(Dropout(0.5))
gru_model.add(GRU(units=32))
gru_model.add(Dropout(0.5))
gru_model.add(Dense(14, activation='softmax'))
gru_model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 526, 300)	5800200
gru (GRU)	(None, 526, 64)	70272
dropout (Dropout)	(None, 526, 64)	0
gru_1 (GRU)	(None, 32)	9408
dropout_1 (Dropout)	(None, 32)	0
dense (Dense)	(None, 14)	462

=====
Total params: 5,880,342
Trainable params: 80,142
Non-trainable params: 5,800,200
=====

```
In [8]: gru_model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy
```

```
In [9]: # Fit model
history_gru=gru_model.fit(X_train,
                           y_train,
                           validation_data=(X_val, y_val),
                           verbose=1,
                           batch_size=256,
                           epochs=30,
                           callbacks=[callback]
                           )
```

```
Epoch 1/30
190/190 [=====] - 21s 61ms/step - loss: 2.2498 - accuracy: 0.33
64 - val_loss: 1.9976 - val_accuracy: 0.4100
Epoch 2/30
190/190 [=====] - 10s 55ms/step - loss: 2.0251 - accuracy: 0.40
78 - val_loss: 1.8992 - val_accuracy: 0.4315
Epoch 3/30
190/190 [=====] - 10s 55ms/step - loss: 1.9526 - accuracy: 0.42
60 - val_loss: 1.8548 - val_accuracy: 0.4426
Epoch 4/30
190/190 [=====] - 11s 55ms/step - loss: 1.9143 - accuracy: 0.43
44 - val_loss: 1.8205 - val_accuracy: 0.4500
Epoch 5/30
190/190 [=====] - 12s 63ms/step - loss: 1.8822 - accuracy: 0.44
15 - val_loss: 1.8025 - val_accuracy: 0.4563
Epoch 6/30
190/190 [=====] - 11s 57ms/step - loss: 1.8621 - accuracy: 0.44
67 - val_loss: 1.7843 - val_accuracy: 0.4607
Epoch 7/30
190/190 [=====] - 11s 56ms/step - loss: 1.8450 - accuracy: 0.44
97 - val_loss: 1.7732 - val_accuracy: 0.4593
```

```

Epoch 8/30
190/190 [=====] - 11s 57ms/step - loss: 1.8236 - accuracy: 0.45
63 - val_loss: 1.7645 - val_accuracy: 0.4622
Epoch 9/30
190/190 [=====] - 11s 57ms/step - loss: 1.8112 - accuracy: 0.45
81 - val_loss: 1.7564 - val_accuracy: 0.4637
Epoch 10/30
190/190 [=====] - 11s 57ms/step - loss: 1.7981 - accuracy: 0.46
39 - val_loss: 1.7595 - val_accuracy: 0.4648
Epoch 11/30
190/190 [=====] - 11s 58ms/step - loss: 1.7848 - accuracy: 0.46
36 - val_loss: 1.7495 - val_accuracy: 0.4641
Epoch 12/30
190/190 [=====] - 11s 58ms/step - loss: 1.7688 - accuracy: 0.46
84 - val_loss: 1.7565 - val_accuracy: 0.4667
Epoch 13/30
190/190 [=====] - 11s 58ms/step - loss: 1.7605 - accuracy: 0.46
91 - val_loss: 1.7496 - val_accuracy: 0.4667
Epoch 14/30
190/190 [=====] - 11s 58ms/step - loss: 1.7468 - accuracy: 0.47
14 - val_loss: 1.7517 - val_accuracy: 0.4711

```

```

In [10]: gru_y_pred = gru_model.predict(X_test)
gru_y_pred_labels = np.argmax(gru_y_pred, axis=1)
gru_y_test_labels = np.argmax(y_test, axis=1)
print(classification_report(gru_y_test_labels, gru_y_pred_labels))

```

```

85/85 [=====] - 2s 17ms/step
              precision    recall  f1-score   support

     0         0.41         0.86         0.56         763
     1         0.43         0.18         0.26         195
     2         0.31         0.10         0.15          90
     3         0.38         0.25         0.30          53
     4         0.24         0.05         0.08         168
     5         0.55         0.47         0.51         278
     6         0.55         0.23         0.32          52
     7         0.78         0.83         0.80         134
     8         0.44         0.22         0.30         206
     9         0.27         0.03         0.06          91
    10         0.54         0.39         0.45         144
    11         0.46         0.25         0.32         252
    12         0.34         0.10         0.16         136
    13         0.53         0.54         0.53         138

 accuracy         0.46         2700
 macro avg        0.45         0.32         0.34         2700
 weighted avg     0.45         0.46         0.41         2700

```

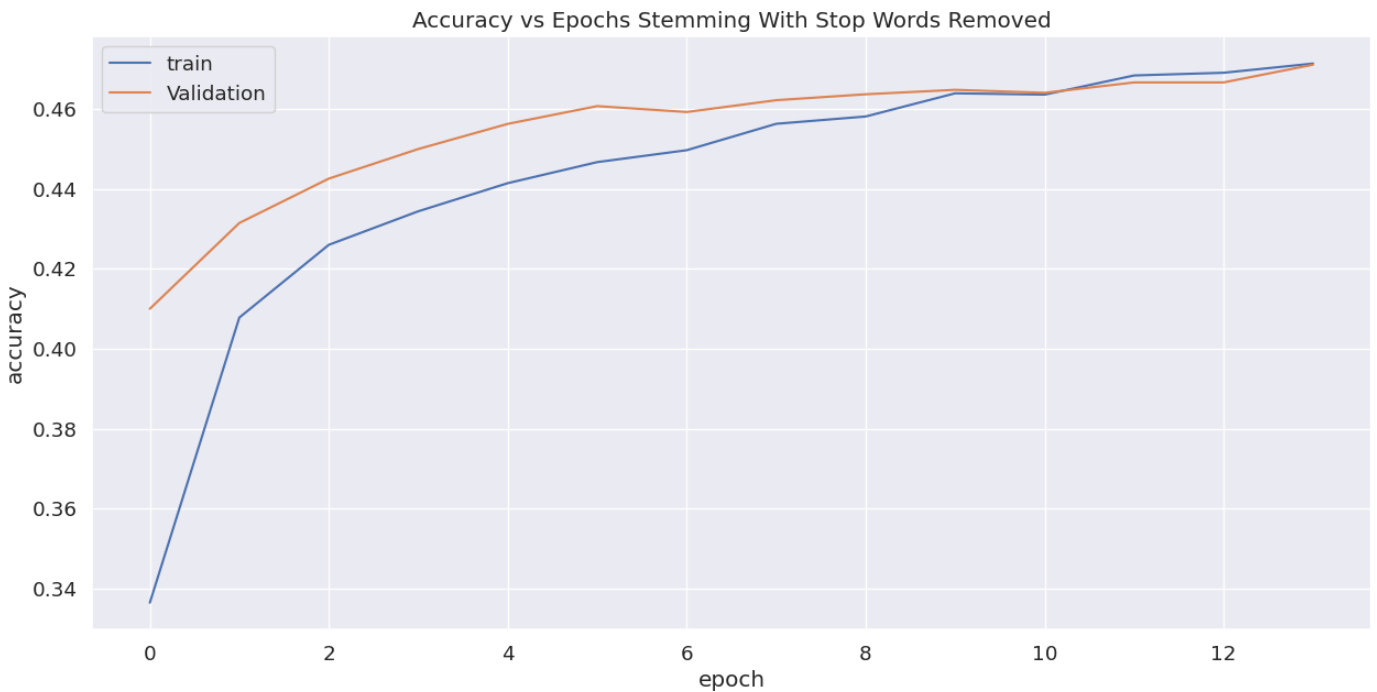
```

In [11]: conf_mat_gru = confusion_matrix(gru_y_test_labels, gru_y_pred_labels)
plt.figure(figsize=(15,7))
# create a heatmap of the confusion matrix
sns.set(font_scale=1.2)
sns.heatmap(conf_mat_gru, annot=True, fmt='d', cmap='Blues', cbar=False,
            xticklabels=['class 0', 'class 1', 'class 2', 'class 3', 'class 4', 'class 5',
                        'class 6', 'class 7', 'class 8', 'class 9', 'class 10',
                        'class 11', 'class 12', 'class 13'],
            yticklabels=['class 0', 'class 1', 'class 2', 'class 3', 'class 4', 'class 5',
                        'class 6', 'class 7', 'class 8', 'class 9', 'class 10',
                        'class 11', 'class 12', 'class 13'])
plt.xlabel('Predicted Labels', fontsize=14)
plt.ylabel('True Labels', fontsize=14)
plt.title('Confusion Matrix Stemming With Stop Words Removed', fontsize=16)
plt.show()

```


Confusion Matrix Stemming With Stop Words Removed														
True Labels	class 0	class 1	class 2	class 3	class 4	class 5	class 6	class 7	class 8	class 9	class 10	class 11	class 12	class 13
	658	17	1	4	7	22	2	0	11	2	1	25	5	8
	120	36	4	2	3	8	1	3	3	0	0	7	5	3
	59	4	9	0	0	6	1	0	3	1	4	2	0	1
	32	1	0	13	1	0	1	0	0	2	1	1	0	1
	115	4	1	1	8	19	0	3	5	0	1	2	1	8
	78	7	3	1	6	132	2	11	9	0	10	4	0	15
	27	1	1	0	1	2	12	0	2	0	1	3	1	1
	7	0	1	0	0	9	0	111	5	0	0	0	0	1
	95	3	1	1	0	18	0	7	46	2	10	6	2	15
	54	2	1	8	1	2	0	0	3	3	5	8	3	1
	58	2	4	0	2	5	0	1	11	0	56	1	0	4
	149	5	1	3	3	6	0	2	5	1	5	62	8	2
	93	0	0	1	0	4	3	3	1	0	2	10	14	5
	41	1	2	0	1	5	0	1	1	0	7	3	2	74
Predicted Labels														

```
In [12]: plt.figure(figsize=(15,7))
plt.plot(history_gru.history['accuracy'])
plt.plot(history_gru.history['val_accuracy'])
plt.title('Accuracy vs Epochs Stemming With Stop Words Removed')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'Validation'], loc='upper left')
plt.show()
```



```
In [13]: plt.figure(figsize=(15,7))
plt.plot(history_gru.history['loss'])
plt.plot(history_gru.history['val_loss'])
plt.title('Loss vs Epochs Stemming With Stop Words Removed')
plt.ylabel('Loss')
plt.xlabel('epoch')
plt.legend(['train', 'Validation'], loc='upper left')
plt.show()
```

Loss vs Epochs Stemming With Stop Words Removed



Stemming With Stop Words Present

```
In [14]: X= df_clean_stopwords_present_stemming.text
y= df_clean_stopwords_present_stemming.labels
X_train, X_other, y_train, y_other = train_test_split(X, y, test_size=0.1, random_state=
X_val,X_test,y_val,y_test= train_test_split(X_other,y_other,test_size=0.5,random_state=7
```

```
In [15]: le = LabelEncoder()
y_train = le.fit_transform(y_train)
y_test = le.transform(y_test)
y_val = le.transform(y_val)

y_train = to_categorical(y_train)
y_test = to_categorical(y_test)
y_val = to_categorical(y_val)
# Tokenize words
tokenizer = Tokenizer(oov_token='UNK')
tokenizer.fit_on_texts(pd.concat([X_train], axis=0))
sequences_train = tokenizer.texts_to_sequences(X_train)
sequences_test = tokenizer.texts_to_sequences(X_test)
sequences_val = tokenizer.texts_to_sequences(X_val)
max_len = max([len(t) for t in X_train])
X_train = pad_sequences(sequences_train, maxlen=max_len, truncating='pre')
X_test = pad_sequences(sequences_test, maxlen=max_len, truncating='pre')
X_val = pad_sequences(sequences_val, maxlen=max_len, truncating='pre')
vocabSize = len(tokenizer.index_word) + 1

#Text Representation Using Glove Embedding
num_tokens = vocabSize
embedding_dim = 300
embeddings_index = {}
misses=0
hits=0
with open(path_to_glove_file) as f:
    for line in f:
        word, coefs = line.split(maxsplit=1)
        coefs = np.fromstring(coefs, "f", sep=" ")
        embeddings_index[word] = coefs
print("Found %s word vectors." % len(embeddings_index))
```

```

embedding_matrix = np.zeros((num_tokens, embedding_dim))
for word, i in tokenizer.word_index.items():
    embedding_vector = embeddings_index.get(word)
    if embedding_vector is not None:
        embedding_matrix[i] = embedding_vector
        hits += 1
    else:
        misses += 1
print("Converted %d words (%d misses)" % (hits, misses))

#define early stopping to control overfitting
callback = EarlyStopping(
    monitor="val_loss",
    patience=3,
    restore_best_weights=True,
)

```

Found 41286 word vectors.
Converted 7014 words (12319 misses)

```

In [16]: gru_model = Sequential()
gru_model.add(Embedding(vocabSize, 300, input_length=X_train.shape[1], weights=[embeddin
gru_model.add(GRU(units=64, return_sequences=True))
gru_model.add(Dropout(0.5))
gru_model.add(GRU(units=32))
gru_model.add(Dropout(0.5))
gru_model.add(Dense(14, activation='softmax'))
gru_model.summary()

```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	(None, 526, 300)	5800200
gru_2 (GRU)	(None, 526, 64)	70272
dropout_2 (Dropout)	(None, 526, 64)	0
gru_3 (GRU)	(None, 32)	9408
dropout_3 (Dropout)	(None, 32)	0
dense_1 (Dense)	(None, 14)	462
=====		
Total params: 5,880,342		
Trainable params: 80,142		
Non-trainable params: 5,800,200		

```

In [17]: gru_model.compile(loss='categorical_crossentropy', optimizer= 'adam', metrics=['accuracy

```

```

In [18]: # Fit model
history_gru=gru_model.fit(X_train,
    y_train,
    validation_data=(X_val, y_val),
    verbose=1,
    batch_size=256,
    epochs=30,
    callbacks=[callback]
)

```

Epoch 1/30
190/190 [=====] - 16s 63ms/step - loss: 2.2322 - accuracy: 0.34


```

20 - val_loss: 1.9711 - val_accuracy: 0.4252
Epoch 2/30
190/190 [=====] - 11s 59ms/step - loss: 2.0015 - accuracy: 0.41
58 - val_loss: 1.8875 - val_accuracy: 0.4437
Epoch 3/30
190/190 [=====] - 11s 59ms/step - loss: 1.9343 - accuracy: 0.43
30 - val_loss: 1.8409 - val_accuracy: 0.4533
Epoch 4/30
190/190 [=====] - 11s 59ms/step - loss: 1.8959 - accuracy: 0.44
40 - val_loss: 1.8103 - val_accuracy: 0.4611
Epoch 5/30
190/190 [=====] - 11s 58ms/step - loss: 1.8645 - accuracy: 0.45
03 - val_loss: 1.7867 - val_accuracy: 0.4611
Epoch 6/30
190/190 [=====] - 11s 58ms/step - loss: 1.8420 - accuracy: 0.45
54 - val_loss: 1.7674 - val_accuracy: 0.4685
Epoch 7/30
190/190 [=====] - 11s 58ms/step - loss: 1.8225 - accuracy: 0.45
96 - val_loss: 1.7566 - val_accuracy: 0.4704
Epoch 8/30
190/190 [=====] - 11s 57ms/step - loss: 1.8038 - accuracy: 0.46
32 - val_loss: 1.7538 - val_accuracy: 0.4678
Epoch 9/30
190/190 [=====] - 11s 58ms/step - loss: 1.7864 - accuracy: 0.46
92 - val_loss: 1.7485 - val_accuracy: 0.4689
Epoch 10/30
190/190 [=====] - 11s 59ms/step - loss: 1.7734 - accuracy: 0.46
97 - val_loss: 1.7361 - val_accuracy: 0.4722
Epoch 11/30
190/190 [=====] - 11s 60ms/step - loss: 1.7607 - accuracy: 0.47
31 - val_loss: 1.7337 - val_accuracy: 0.4726
Epoch 12/30
190/190 [=====] - 11s 60ms/step - loss: 1.7480 - accuracy: 0.47
56 - val_loss: 1.7272 - val_accuracy: 0.4759
Epoch 13/30
190/190 [=====] - 11s 60ms/step - loss: 1.7361 - accuracy: 0.47
98 - val_loss: 1.7410 - val_accuracy: 0.4670
Epoch 14/30
190/190 [=====] - 11s 58ms/step - loss: 1.7268 - accuracy: 0.48
14 - val_loss: 1.7313 - val_accuracy: 0.4752
Epoch 15/30
190/190 [=====] - 11s 57ms/step - loss: 1.7146 - accuracy: 0.48
54 - val_loss: 1.7258 - val_accuracy: 0.4785
Epoch 16/30
190/190 [=====] - 11s 57ms/step - loss: 1.7038 - accuracy: 0.48
77 - val_loss: 1.7296 - val_accuracy: 0.4744
Epoch 17/30
190/190 [=====] - 11s 57ms/step - loss: 1.6901 - accuracy: 0.49
04 - val_loss: 1.7296 - val_accuracy: 0.4737
Epoch 18/30
190/190 [=====] - 11s 57ms/step - loss: 1.6787 - accuracy: 0.49
27 - val_loss: 1.7308 - val_accuracy: 0.4737

```

```

In [19]: gru_y_pred = gru_model.predict(X_test)
gru_y_pred_labels = np.argmax(gru_y_pred, axis=1)
gru_y_test_labels = np.argmax(y_test, axis=1)
print(classification_report(gru_y_test_labels, gru_y_pred_labels))

```

```

85/85 [=====] - 2s 11ms/step
              precision    recall  f1-score   support

     0         0.42         0.82         0.56         763
     1         0.43         0.22         0.29         195
     2         0.31         0.12         0.17          90
     3         0.39         0.26         0.31          53
     4         0.23         0.04         0.07         168

```

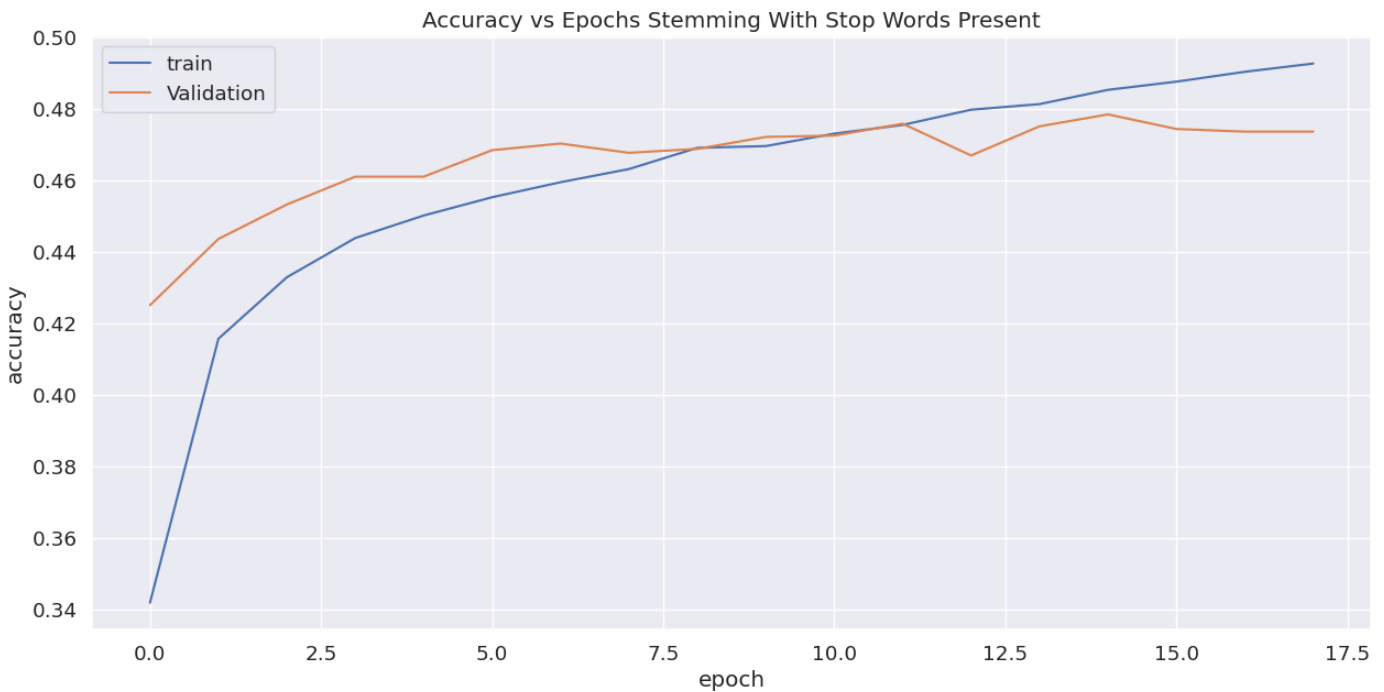
5	0.52	0.47	0.50	278
6	0.59	0.25	0.35	52
7	0.77	0.83	0.80	134
8	0.47	0.18	0.26	206
9	0.33	0.12	0.18	91
10	0.50	0.41	0.45	144
11	0.50	0.36	0.42	252
12	0.33	0.15	0.20	136
13	0.53	0.53	0.53	138

accuracy			0.46	2700
macro avg	0.45	0.34	0.36	2700
weighted avg	0.45	0.46	0.42	2700

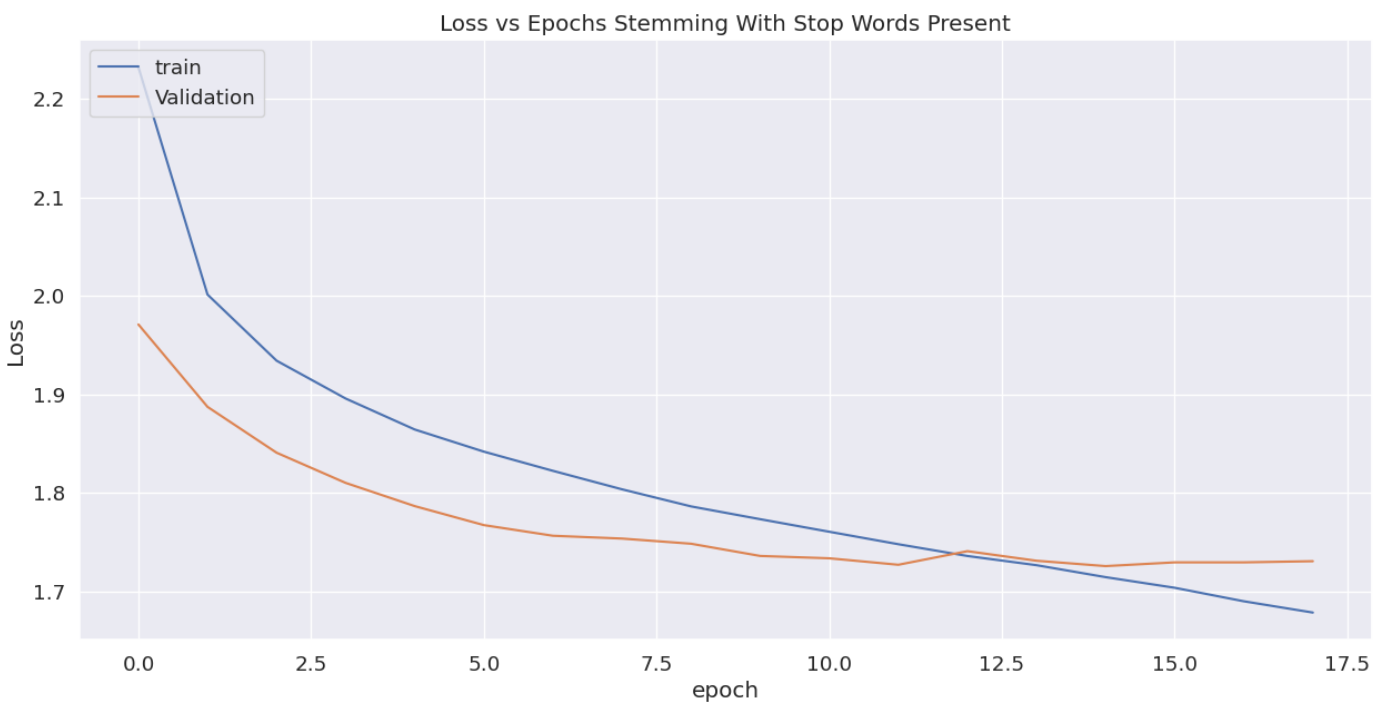
```
In [20]: conf_mat_gru = confusion_matrix(gru_y_test_labels, gru_y_pred_labels)
plt.figure(figsize=(15,7))
# create a heatmap of the confusion matrix
sns.set(font_scale=1.2)
sns.heatmap(conf_mat_gru, annot=True, fmt='d', cmap='Blues', cbar=False,
            xticklabels=['class 0', 'class 1', 'class 2', 'class 3', 'class 4', 'class 5',
                        'class 6', 'class 7', 'class 8', 'class 9', 'class 10',
                        'class 11', 'class 12', 'class 13'],
            yticklabels=['class 0', 'class 1', 'class 2', 'class 3', 'class 4', 'class 5',
                        'class 6', 'class 7', 'class 8', 'class 9', 'class 10',
                        'class 11', 'class 12', 'class 13'])
plt.xlabel('Predicted Labels', fontsize=14)
plt.ylabel('True Labels', fontsize=14)
plt.title('Confusion Matrix Stemming With Stop Words Present', fontsize=16)
plt.show()
```

Confusion Matrix Stemming With Stop Words Present															
True Labels	class 0	622	20	4	6	7	31	2	0	9	6	7	28	11	10
	class 1	107	43	4	1	3	12	1	3	3	3	0	4	7	4
	class 2	58	3	11	0	2	5	1	0	0	1	3	3	1	2
	class 3	27	2	0	14	1	0	1	0	0	3	1	3	0	1
	class 4	109	7	0	1	7	21	0	3	4	1	4	3	1	7
	class 5	80	7	5	1	5	131	2	11	6	1	12	3	1	13
	class 6	25	0	2	0	0	2	13	0	2	0	1	5	1	1
	class 7	7	0	0	1	0	9	0	111	3	0	1	1	0	1
	class 8	92	6	2	2	0	15	0	8	38	3	13	9	3	15
	class 9	49	2	0	7	2	2	0	0	1	11	4	8	4	1
	class 10	54	1	5	0	0	5	0	2	11	1	59	3	1	2
	class 11	123	6	1	2	2	6	0	2	3	2	4	90	8	3
	class 12	80	3	0	1	0	3	2	3	1	1	1	17	20	4
	class 13	41	0	2	0	1	8	0	1	0	0	8	2	2	73
			class 0	class 1	class 2	class 3	class 4	class 5	class 6	class 7	class 8	class 9	class 10	class 11	class 12
Predicted Labels															

```
In [21]: plt.figure(figsize=(15,7))
plt.plot(history_gru.history['accuracy'])
plt.plot(history_gru.history['val_accuracy'])
plt.title('Accuracy vs Epochs Stemming With Stop Words Present')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'Validation'], loc='upper left')
plt.show()
```



```
In [22]: plt.figure(figsize=(15,7))
plt.plot(history_gru.history['loss'])
plt.plot(history_gru.history['val_loss'])
plt.title('Loss vs Epochs Stemming With Stop Words Present')
plt.ylabel('Loss')
plt.xlabel('epoch')
plt.legend(['train', 'Validation'], loc='upper left')
plt.show()
```



Lemmatization With Stop Words Present

```
In [23]: X= df_clean_stopwords_present_lemma.text
y= df_clean_stopwords_present_lemma.labels
X_train, X_other, y_train, y_other = train_test_split(X, y, test_size=0.1, random_state=7)
X_val,X_test,y_val,y_test= train_test_split(X_other,y_other,test_size=0.5,random_state=7)
```

```
In [24]: le = LabelEncoder()
```

```

y_train = le.fit_transform(y_train)
y_test = le.transform(y_test)
y_val = le.transform(y_val)

y_train = to_categorical(y_train)
y_test = to_categorical(y_test)
y_val = to_categorical(y_val)
# Tokenize words
tokenizer = Tokenizer(oov_token='UNK')
tokenizer.fit_on_texts(pd.concat([X_train], axis=0))
sequences_train = tokenizer.texts_to_sequences(X_train)
sequences_test = tokenizer.texts_to_sequences(X_test)
sequences_val = tokenizer.texts_to_sequences(X_val)
max_len = max([len(t) for t in X_train])
X_train = pad_sequences(sequences_train, maxlen=max_len, truncating='pre')
X_test = pad_sequences(sequences_test, maxlen=max_len, truncating='pre')
X_val = pad_sequences(sequences_val, maxlen=max_len, truncating='pre')
vocabSize = len(tokenizer.index_word) + 1

#Text Representation Using Glove Embedding
num_tokens = vocabSize
embedding_dim = 300
embeddings_index = {}
misses=0
hits=0
with open(path_to_glove_file) as f:
    for line in f:
        word, coefs = line.split(maxsplit=1)
        coefs = np.fromstring(coefs, "f", sep=" ")
        embeddings_index[word] = coefs
print("Found %s word vectors." % len(embeddings_index))

embedding_matrix = np.zeros((num_tokens, embedding_dim))
for word, i in tokenizer.word_index.items():
    embedding_vector = embeddings_index.get(word)
    if embedding_vector is not None:
        embedding_matrix[i] = embedding_vector
        hits += 1
    else:
        misses += 1
print("Converted %d words (%d misses)" % (hits, misses))

#define early stopping to control overfitting
callback = EarlyStopping(
    monitor="val_loss",
    patience=3,
    restore_best_weights=True,
)

```

Found 59131 word vectors.
Converted 7936 words (11397 misses)

```

In [25]: gru_model = Sequential()
gru_model.add(Embedding(vocabSize, 300, input_length=X_train.shape[1], weights=[embeddin
gru_model.add(GRU(units=64, return_sequences=True))
gru_model.add(Dropout(0.5))
gru_model.add(GRU(units=32))
gru_model.add(Dropout(0.5))
gru_model.add(Dense(14, activation='softmax'))
gru_model.summary()

```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
=====		

embedding_2 (Embedding)	(None, 526, 300)	5800200
gru_4 (GRU)	(None, 526, 64)	70272
dropout_4 (Dropout)	(None, 526, 64)	0
gru_5 (GRU)	(None, 32)	9408
dropout_5 (Dropout)	(None, 32)	0
dense_2 (Dense)	(None, 14)	462

=====

Total params: 5,880,342
Trainable params: 80,142
Non-trainable params: 5,800,200

```
In [26]: gru_model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy
```

```
In [27]: # Fit model
history_gru=gru_model.fit(X_train,
                           y_train,
                           validation_data=(X_val, y_val),
                           verbose=1,
                           batch_size=256,
                           epochs=30,
                           callbacks=[callback]
                           )
```

```
Epoch 1/30
190/190 [=====] - 16s 61ms/step - loss: 2.2397 - accuracy: 0.33
60 - val_loss: 1.9760 - val_accuracy: 0.4070
Epoch 2/30
190/190 [=====] - 11s 58ms/step - loss: 2.0059 - accuracy: 0.41
09 - val_loss: 1.8751 - val_accuracy: 0.4456
Epoch 3/30
190/190 [=====] - 11s 58ms/step - loss: 1.9260 - accuracy: 0.43
40 - val_loss: 1.8204 - val_accuracy: 0.4559
Epoch 4/30
190/190 [=====] - 11s 58ms/step - loss: 1.8873 - accuracy: 0.44
34 - val_loss: 1.7921 - val_accuracy: 0.4581
Epoch 5/30
190/190 [=====] - 11s 57ms/step - loss: 1.8543 - accuracy: 0.45
04 - val_loss: 1.7654 - val_accuracy: 0.4667
Epoch 6/30
190/190 [=====] - 11s 58ms/step - loss: 1.8298 - accuracy: 0.45
90 - val_loss: 1.7473 - val_accuracy: 0.4689
Epoch 7/30
190/190 [=====] - 11s 57ms/step - loss: 1.8123 - accuracy: 0.46
56 - val_loss: 1.7426 - val_accuracy: 0.4730
Epoch 8/30
190/190 [=====] - 11s 57ms/step - loss: 1.7900 - accuracy: 0.46
90 - val_loss: 1.7332 - val_accuracy: 0.4711
Epoch 9/30
190/190 [=====] - 11s 58ms/step - loss: 1.7726 - accuracy: 0.47
27 - val_loss: 1.7261 - val_accuracy: 0.4748
Epoch 10/30
190/190 [=====] - 11s 58ms/step - loss: 1.7582 - accuracy: 0.47
67 - val_loss: 1.7153 - val_accuracy: 0.4737
Epoch 11/30
190/190 [=====] - 11s 57ms/step - loss: 1.7449 - accuracy: 0.48
02 - val_loss: 1.7160 - val_accuracy: 0.4737
Epoch 12/30
190/190 [=====] - 11s 58ms/step - loss: 1.7327 - accuracy: 0.48
```

```

05 - val_loss: 1.7082 - val_accuracy: 0.4733
Epoch 13/30
190/190 [=====] - 11s 57ms/step - loss: 1.7204 - accuracy: 0.48
59 - val_loss: 1.7174 - val_accuracy: 0.4807
Epoch 14/30
190/190 [=====] - 11s 57ms/step - loss: 1.7086 - accuracy: 0.48
79 - val_loss: 1.7040 - val_accuracy: 0.4811
Epoch 15/30
190/190 [=====] - 11s 58ms/step - loss: 1.6982 - accuracy: 0.48
99 - val_loss: 1.7085 - val_accuracy: 0.4826
Epoch 16/30
190/190 [=====] - 11s 58ms/step - loss: 1.6867 - accuracy: 0.49
28 - val_loss: 1.7160 - val_accuracy: 0.4819
Epoch 17/30
190/190 [=====] - 11s 58ms/step - loss: 1.6732 - accuracy: 0.49
63 - val_loss: 1.7111 - val_accuracy: 0.4815

```

```

In [28]: gru_y_pred = gru_model.predict(X_test)
gru_y_pred_labels = np.argmax(gru_y_pred, axis=1)
gru_y_test_labels = np.argmax(y_test, axis=1)
print(classification_report(gru_y_test_labels, gru_y_pred_labels))

```

```

85/85 [=====] - 2s 13ms/step
              precision    recall  f1-score   support

     0         0.43         0.82         0.57         763
     1         0.49         0.19         0.28         195
     2         0.28         0.12         0.17          90
     3         0.43         0.36         0.39          53
     4         0.41         0.15         0.22         168
     5         0.54         0.46         0.50         278
     6         0.57         0.23         0.33          52
     7         0.75         0.84         0.79         134
     8         0.48         0.21         0.30         206
     9         0.25         0.07         0.10          91
    10         0.52         0.35         0.42         144
    11         0.52         0.40         0.45         252
    12         0.37         0.13         0.19         136
    13         0.52         0.54         0.53         138

 accuracy                   0.47         2700
 macro avg                 0.47         0.35         0.37         2700
 weighted avg              0.47         0.47         0.43         2700

```

```


In [29]: conf_mat_gru = confusion_matrix(gru_y_test_labels, gru_y_pred_labels)
plt.figure(figsize=(15,7))
# create a heatmap of the confusion matrix
sns.set(font_scale=1.2)
sns.heatmap(conf_mat_gru, annot=True, fmt='d', cmap='Blues', cbar=False,
            xticklabels=['class 0', 'class 1', 'class 2', 'class 3', 'class 4', 'class 5',
                        'class 6', 'class 7', 'class 8', 'class 9', 'class 10',
                        'class 11', 'class 12', 'class 13'],
            yticklabels=['class 0', 'class 1', 'class 2', 'class 3', 'class 4', 'class 5',
                        'class 6', 'class 7', 'class 8', 'class 9', 'class 10',
                        'class 11', 'class 12', 'class 13'])
plt.xlabel('Predicted Labels', fontsize=14)
plt.ylabel('True Labels', fontsize=14)
plt.title('Confusion Matrix Lemmatization With Stop Words Present', fontsize=16)
plt.show()

```

Confusion Matrix Lemmatization With Stop Words Present

True Labels \ Predicted Labels	class 0	class 1	class 2	class 3	class 4	class 5	class 6	class 7	class 8	class 9	class 10	class 11	class 12	class 13
class 0	629	12	5	5	12	27	2	1	11	4	3	32	7	13
class 1	113	38	4	3	7	10	0	3	4	1	0	5	4	3
class 2	59	4	11	0	1	5	0	0	0	3	2	3	1	1
class 3	24	1	0	19	1	0	0	0	0	1	2	3	0	2
class 4	97	4	1	1	25	19	0	3	3	1	1	3	1	9
class 5	85	6	2	1	7	129	2	12	7	0	11	3	1	12
class 6	24	0	2	0	0	2	12	0	3	0	1	5	2	1
class 7	7	0	0	1	0	8	0	112	3	0	1	1	0	1
class 8	89	4	4	2	1	16	1	9	44	2	9	7	4	14
class 9	45	2	3	8	2	2	1	0	0	6	6	12	3	1
class 10	60	1	5	0	1	2	0	3	11	3	51	2	0	5
class 11	112	4	1	3	2	6	1	2	4	2	4	102	6	3
class 12	79	1	0	1	1	4	2	3	2	1	1	18	18	5
class 13	40	0	2	0	1	8	0	1	0	0	7	2	2	75

```
In [31]: plt.figure(figsize=(15,7))
plt.plot(history_gru.history['accuracy'])
plt.plot(history_gru.history['val_accuracy'])
plt.title('Accuracy vs Epoch Lemmatization With Stop Words Present')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'Validation'], loc='upper left')
plt.show()
```



```
In [32]: plt.figure(figsize=(15,7))
plt.plot(history_gru.history['loss'])
plt.plot(history_gru.history['val_loss'])
plt.title('Loss vs Epoch Lemmatization With Stop Words Present')
plt.ylabel('Loss')
plt.xlabel('epoch')
plt.legend(['train', 'Validation'], loc='upper left')
plt.show()
```

Loss vs Epoch Lemmatization With Stop Words Present

